

CHAPTER 4

Entities and Relationships

Abstract

This chapter discusses data modeling independent of any specific theoretical data model. Topics include entities, attributes, using entity-relationship diagrams to represent entities and their attributes, domains, data relationships, representing data relationships in entity-relationship diagrams (three styles of diagramming), and handling many-to-many relationships using composite entities.

Keywords

entity
entity relationship diagram
database column domains
database attributes
database relationships
one-to-one relationships
one-to-many relationships
many-to-many relationships
weak entities
mandatory relationships
composite entities
data modeling
data flow

In this chapter, we will look at the basis for all database systems: the relationships between elements in the database environment. The formal way in which you express data relationships to a database management system (DBMS) is known as a *data model*. The relational data model, about which you will learn in this book, is just such a formal structure. However, the underlying relationships in a database are independent of the data model and, therefore, also independent of the DBMS you are using. Before you can design a database for any data model, you need to be able to identify data relationships.

Note: Most DBMSs support only one data model. Therefore, when you choose a DBMS, you are also choosing your data model.

Entities and Their Attributes

Most database designs begin by identifying entities. An *entity* is something about which we store data. A customer is an entity, as is a merchandise item stocked by *Antique Opticals*. Entities are not necessarily tangible. For example, an event such as a concert is an entity; an appointment to see a doctor is an entity.

Entities have data that describe them (their *attributes*). For example, a customer entity is usually described by a customer number, first name, last name, street, city, state, zip code, and phone number. A concert entity might be described by a title, date, location, and name of the performer.

When we represent entities in a database, we actually store only the attributes. Each group of attributes that describes a single real-world occurrence of an entity acts to represent an *instance* of an entity. For example, in [Figure 4.1](#), you can see four instances of a customer entity stored in a database. If we have 1000 customers in our database, then there will be 1000 collections of customer attributes.



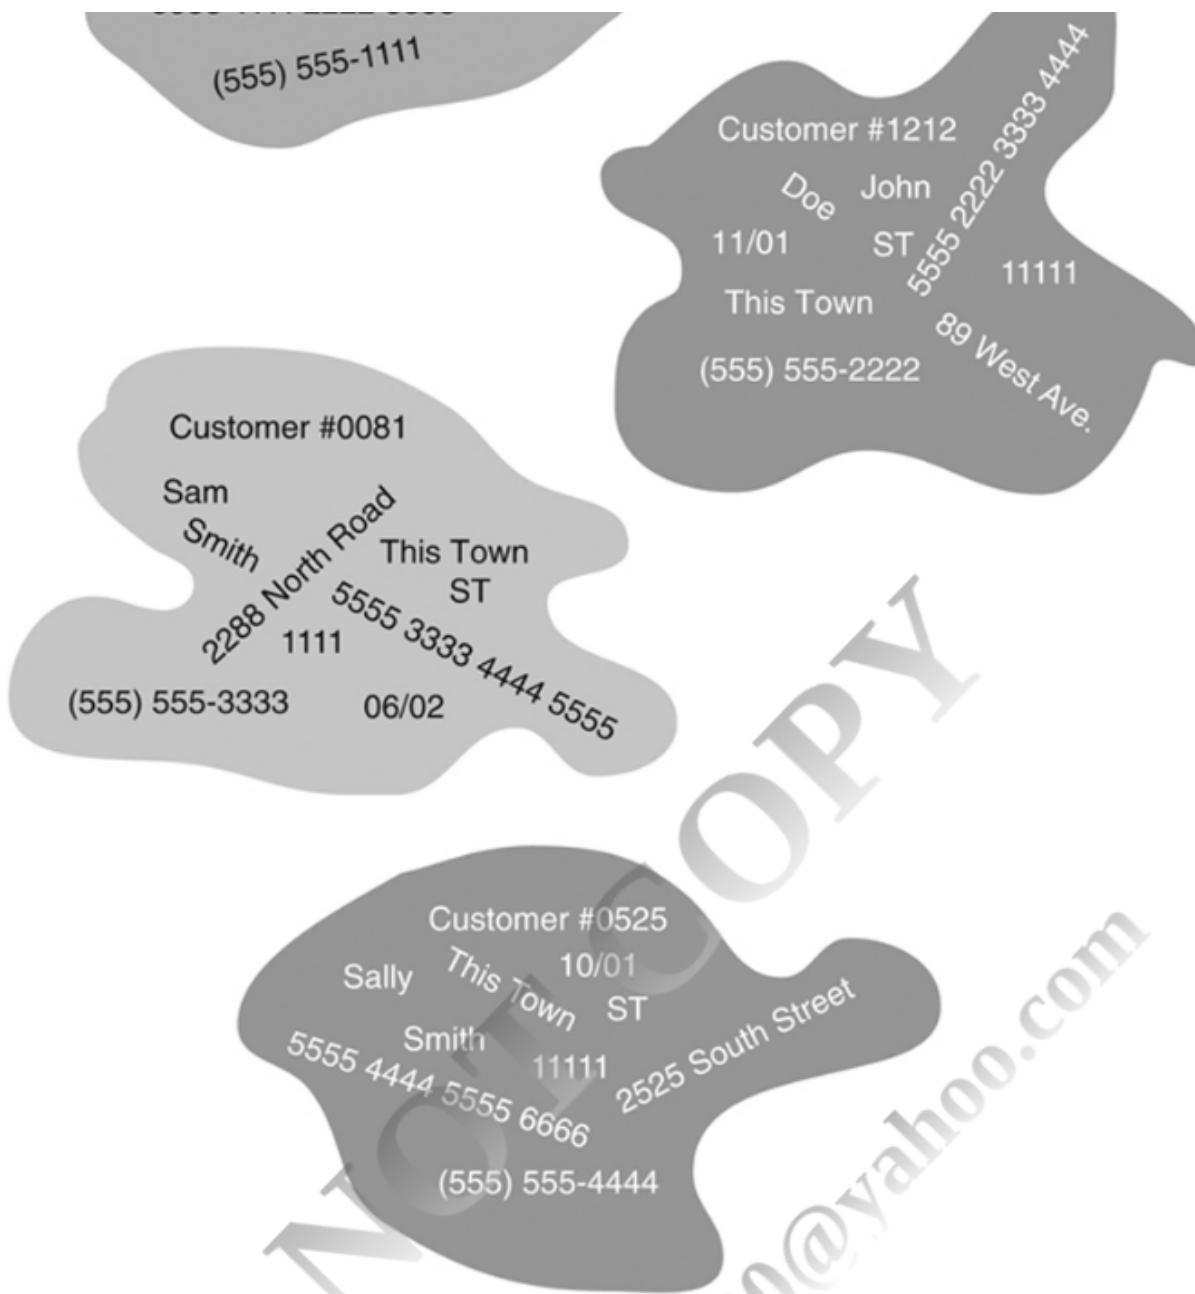


FIGURE 4.1 Instances of a customer entity in a database.

Note: Keep in mind that we are not making any statements about how the instances are physically stored. What you see in Figure 4.1 is purely a conceptual representation.

Entity Identifiers

The only purpose for putting the data that describe an entity into a database (an instance of the entity) is to retrieve the data at some later date. This means that we must have some way of distinguishing one entity instance from another, so that we can always be certain that we are retrieving the precise instance we want. We do this by ensuring that each entity instance has some attribute values that distinguish it from every other instance of the same entity in the database (an *entity identifier*).

Assume, for example, that *Antique Opticals* has only two customers named John Smith. If an employee searches for the items John Smith has on order, which John Smith will the DBMS retrieve? In this case, both of them. Because there is no way to distinguish between the two customers, the result of the query will be inaccurate.

Antique Opticals solved the problem by creating unique customer numbers. That is indeed a common solution to identifying instances of entities where there is no simple unique identifier suggested by the data themselves.

Another solution would be to pair the customer's first and last names with his or her telephone number. This combination of pieces of data (a *concatenated identifier*) would also uniquely identify each customer. There are, however, two drawbacks to doing so. First, the identifier is long and clumsy; it would be easy to make mistakes when entering any of the parts. Second, if the customer's phone number changes, then the identifier must also change. As you read in Chapter 3, changes in an entity identifier can cause serious problems in a database.

Some entities, such as invoices, come with *natural identifiers* (the invoice number). We assign unique, meaningless numbers to others—especially accounts, people, places, and things. Still others require concatenated identifiers.

Note: We will examine the issue of what makes a good unique identifier more closely in Chapter 5, when we talk about “primary keys.”

Printed by: rituadhirakari20@yahoo.com. Printing is for personal, private use only. No part of this book may be reproduced or transmitted without publisher's prior permission. Violators will be prosecuted.

When we store an instance of an entity in a database, we want the DBMS to ensure that the new instance has a unique identifier. This is an example of a *constraint* on a database, a rule to which data must adhere. The enforcement of a variety of database constraints helps us to maintain data consistency and accuracy.

Single-Valued Versus Multivalued Attributes

Because we are eventually going to create a relational database, the attributes in our data model must be *single-valued*. This means that for a given instance of an entity, each attribute can have only one value. For example, the customer entity shown in [Figure 4.1](#) allows only one telephone number for each customer. If a customer has more than one phone number, and wants them all included in the database, then the customer entity cannot handle them.

Note: While it is true that the conceptual data model of a database is independent of the formal data model used to express the structure of the data to a DBMS, we often make decisions on how to model the data based on the requirements of the formal data model we will be using. Removing multivalued attributes is one such case. You will also see an example of this when we deal with many-to-many relationships between entities, later in this chapter.

The existence of more than one phone number turns the phone number attribute into a *multivalued attribute*. Because an entity in a relational database cannot have multivalued attributes, you must handle those attributes by creating an entity to hold them.

In the case of the multiple phone numbers, we could create a phone number entity. Each instance of the entity would include the customer number of the person to whom the phone number belonged, along with the telephone number. If a customer had three phone numbers, then there would be three instances of the phone number entity for the customer. The entity's identifier would be the concatenation of the customer number and the telephone number.

Note: There is no way to avoid using the telephone number as part of the entity identifier in the telephone number entity. As you will come to understand as you read this book, in this particular case, there is no harm in using it in this way.

Note: Some people view a telephone number as made of three distinct pieces of data: an area code, an exchange, and a unique number. However, in common use, we generally consider a telephone number to be a single value.

What is the problem with multivalued attributes? Multivalued attributes can cause problems with the meaning of data in the database, significantly slow down searching, and place unnecessary restrictions on the amount of data that can be stored.

Assume, for example, that you have an Employee entity, with attributes for the name and birthdates of dependents. Each attribute is allowed to store multiple values, as in [Figure 4.2](#), where each gray blob represents a single instance of the Employee entity. How will you associate the correct birthdate with the name of the dependent to which it applies? Will it be by the position of a value stored in the attribute (in other words, the first name is related to the first birthdate, and so on)? If so, how will you ensure that there is a birthdate for each name, and a name for each birthdate? How will you ensure that the order of the values is never mixed up?

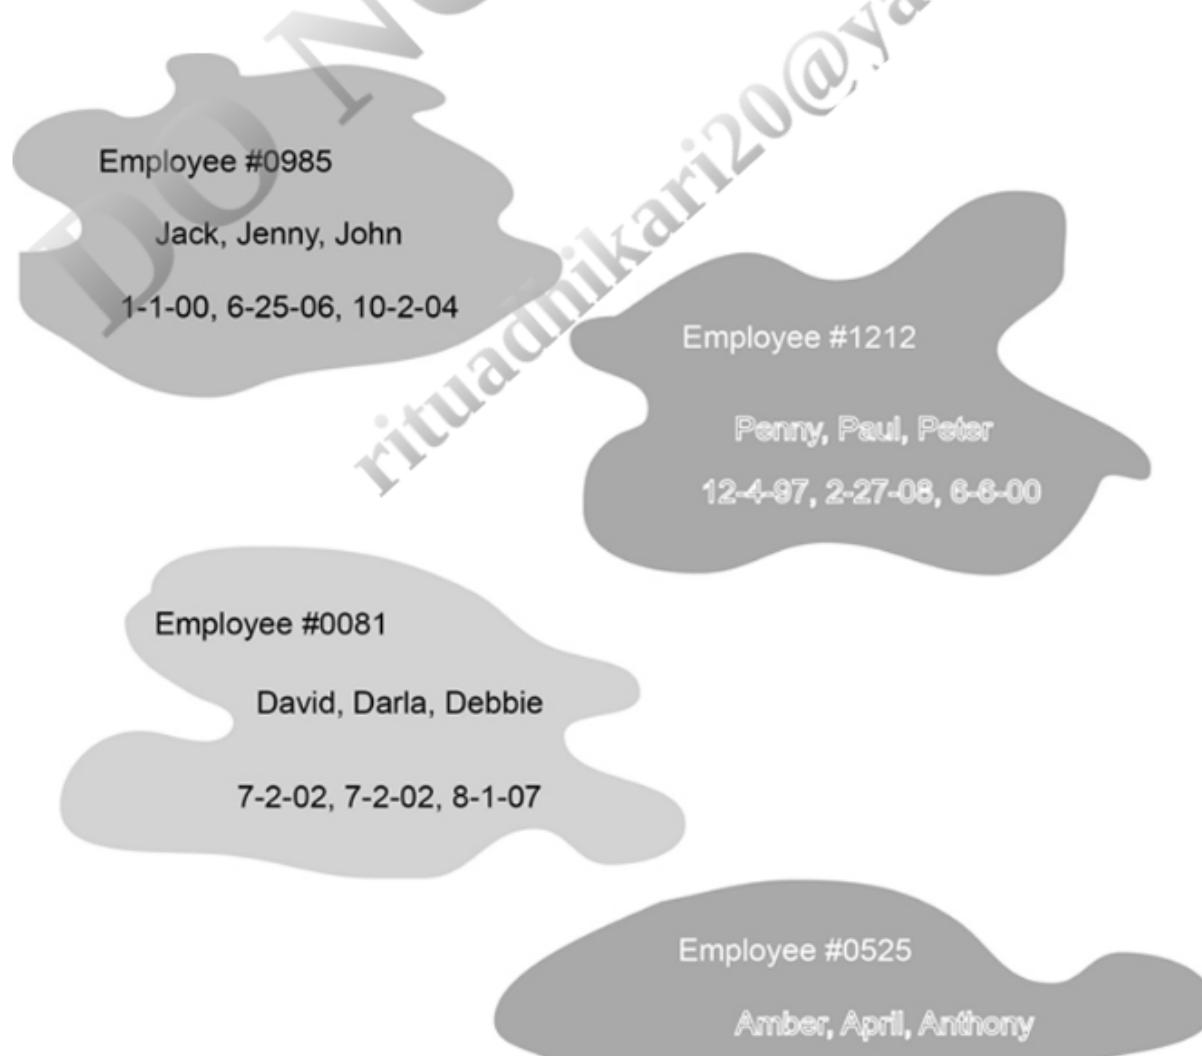


FIGURE 4.2 Entity instances containing multivalued attributes.

When searching a multivalued attribute, a DBMS must search each value in the attribute, most likely scanning the contents of the attribute sequentially. A sequential search is the slowest type of search available.

In addition, how many values should a multivalued attribute be able to store? If you specify a maximum number, what will happen when you need to store more than the maximum number of values? For example, what if you allow room for 10 dependents in the Employee entity just discussed, and you encounter an employee with 11 dependents? Do you create another instance of the Employee entity for that person? Consider all the problems that doing so would create, particularly in terms of the unnecessary duplicated data.

Note: Although it is theoretically possible to write a DBMS that will store an unlimited number of values in an attribute, the implementation would be difficult, and searching much slower than if the maximum number of values were specified in the database design.

As a general rule, if you run across a multivalued attribute, this is a major hint that you need another entity. The only way to handle multiple values of the same attribute is to create an entity of which you can store multiple instances, one for each value of the attribute (for example, [Figure 4.3](#)). In the case of the Employee entity, we would need a Dependent entity that could be related to the Employee entity. There would be one instance of the Dependent entity related to an instance of the Employee entity, for each of an employee's dependents. In this way, there is no limit to the number of an employee's dependents. In addition, each instance of the Dependent entity would contain the name and birthdate of only one dependent, eliminating any confusion about which name was associated with which birthdate. Searching would also be faster, because the DBMS could use fast searching techniques on the individual Dependent entity instances, without resorting to the slow sequential search.

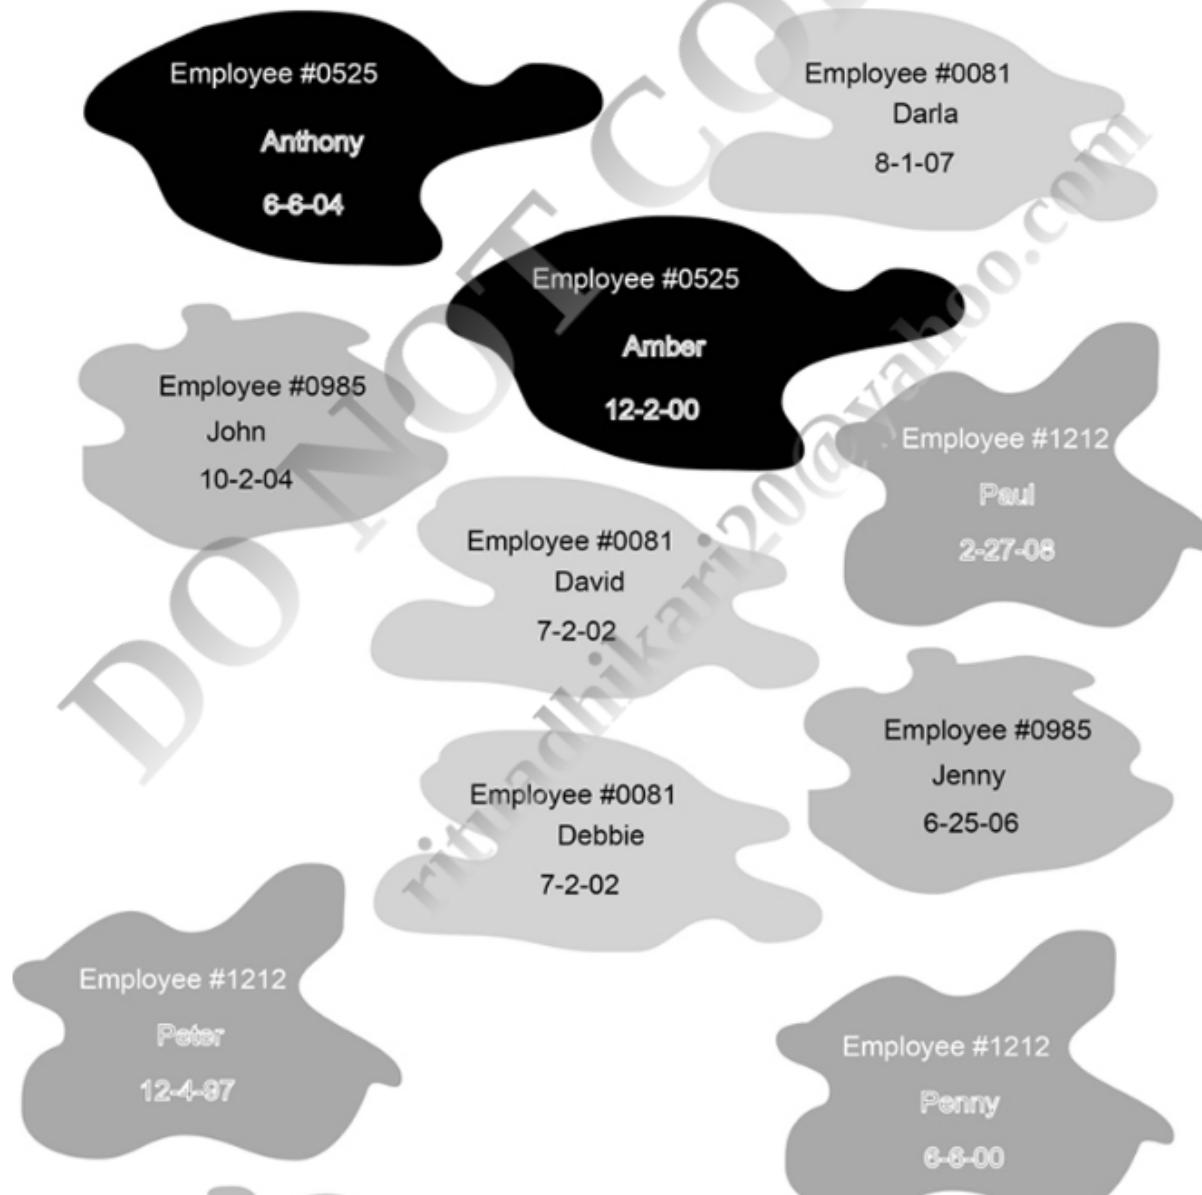




FIGURE 4.3 Using multiple instances of an entity to handle a multivalued attribute.

Avoiding Collections of Entities

When you first begin to work with entities, the nature of an entity can be somewhat confusing. Consider, for example, the merchandise inventory handled by *Antique Opticals*. Is “inventory” an entity? No. Inventory is a collection of the merchandise items handled by the store. The entity is actually the merchandise item. Viewing all of the instances of the merchandise item entity as a whole provides the inventory.

To make this a bit clearer, consider the attributes you would need if you decided to include an inventory entity: merchandise item number, item title, number in stock, retail price, and so on. But because you are trying to describe an entire inventory with a single entity, you need multiple values for each of those attributes. As you read earlier, however, attributes should not be multivalued. This tells you that inventory cannot stand as an entity. It must be represented as a collection of instances of a merchandise item entity.

As another example, consider a person’s medical history maintained by a doctor. Like an inventory, a medical history is a collection of more than one entity. A medical history is made up of appointments and the events that occur during those appointments. Therefore, the history is really a collection of instances of appointment entities and medical treatment entities. The “history” is an output that a database application can obtain by gathering the data stored in the underlying instances.

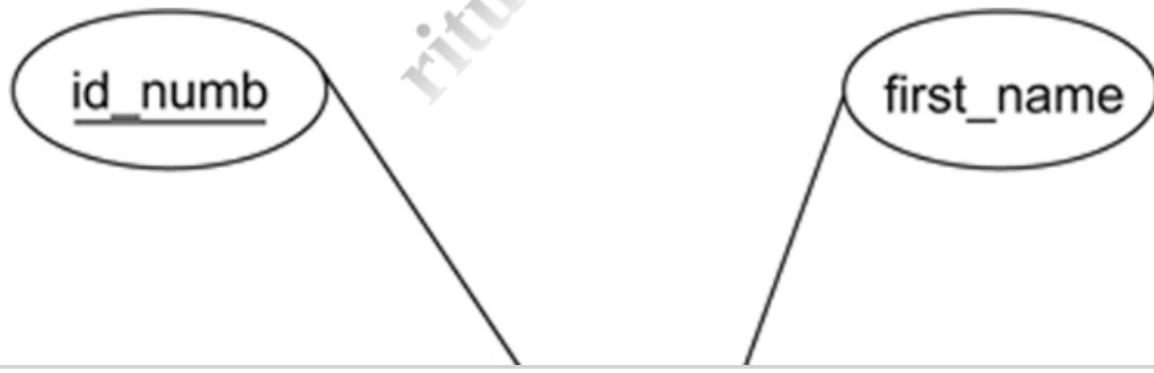
Documenting Entities and Their Attributes

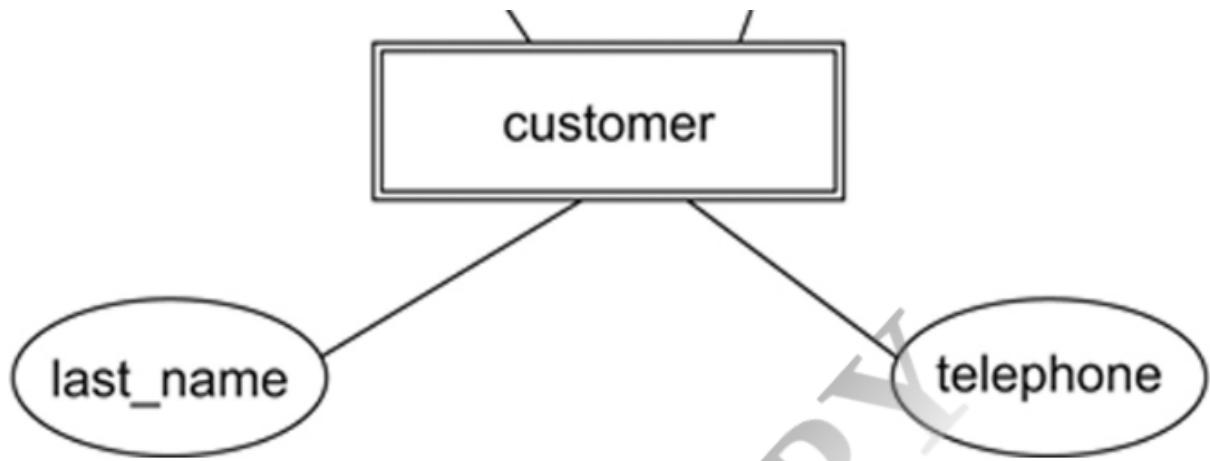
Entity-relationship (ER) diagrams (ERDs) provide a way to document the entities in a database, along with the attributes that describe them. There are actually several styles of ER diagrams. Today, there are three major methods: the Chen model (named after the originator of ER modeling, Dr. Peter P.S. Chen), Information Engineering (IE, or “crows feet”), and Unified Modeling Language (UML). If you are not including object-oriented concepts in a data model, it really doesn’t matter which you use, as long as everyone who is using the diagram understands the symbols. However, UML is specifically intended for the object-oriented environment, and is usually the choice when objects are included.

All three diagramming styles use rectangles to represent entities. Each entity’s name appears in the rectangle, and is expressed in the singular, as in



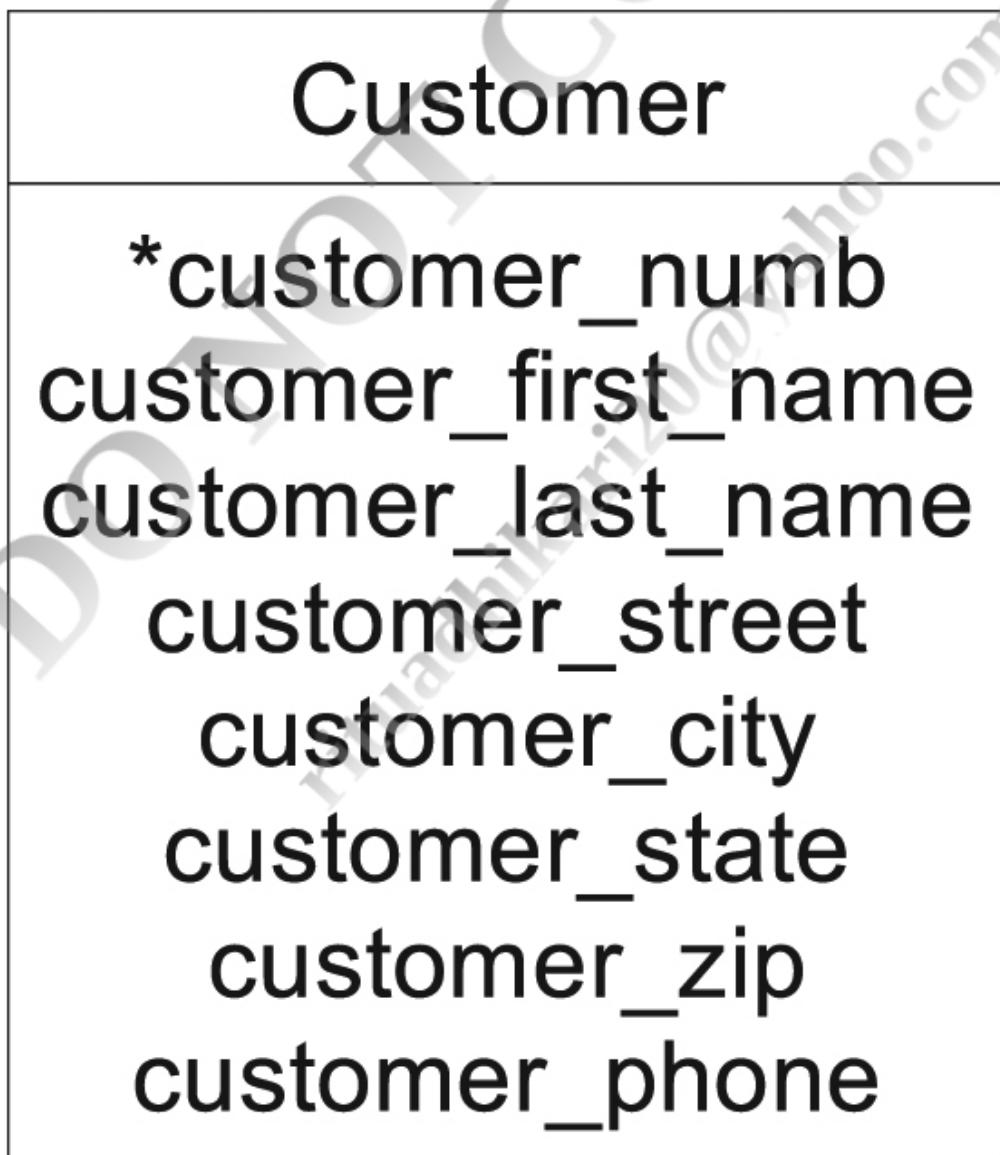
The original Chen model has no provision for showing attributes on the ER diagram itself. However, many people have extended the model to include the attributes in ovals:





The entity's identifier is underlined (id_numb).

The IE and UML styles of ER diagramming include the attributes in the rectangle with the entity; the entity identifier is preceded by an asterisk (*):



Because the IE and UML approaches tend to produce a less cluttered diagram, and because they are the more flexible styles, we will be using IE for most of the diagrams in this book, although you will be introduced to elements of the Chen style and the UML style throughout this chapter. When we introduce object-oriented concepts, we will be using the UML diagramming style.

Printed by: rituadhibkari20@yahoo.com. Printing is for personal, private use only. No part of this book may be reproduced or transmitted without publisher's prior permission. Violators will be prosecuted.

Entities and Attributes for *Antique Opticals*

The major entities and their attributes for the *Antique Opticals* database can be found in [Figure 4.4](#). As you will see, the design will require additional entities as we work with the relationships between those already identified. In particular, there is no information in [Figure 4.4](#) that indicates which items appear on which orders, and no information about which used laser discs are purchased by the store during a single transaction. This occurs because the missing information is a part of the logical relationships between customers, orders, purchases, and items.

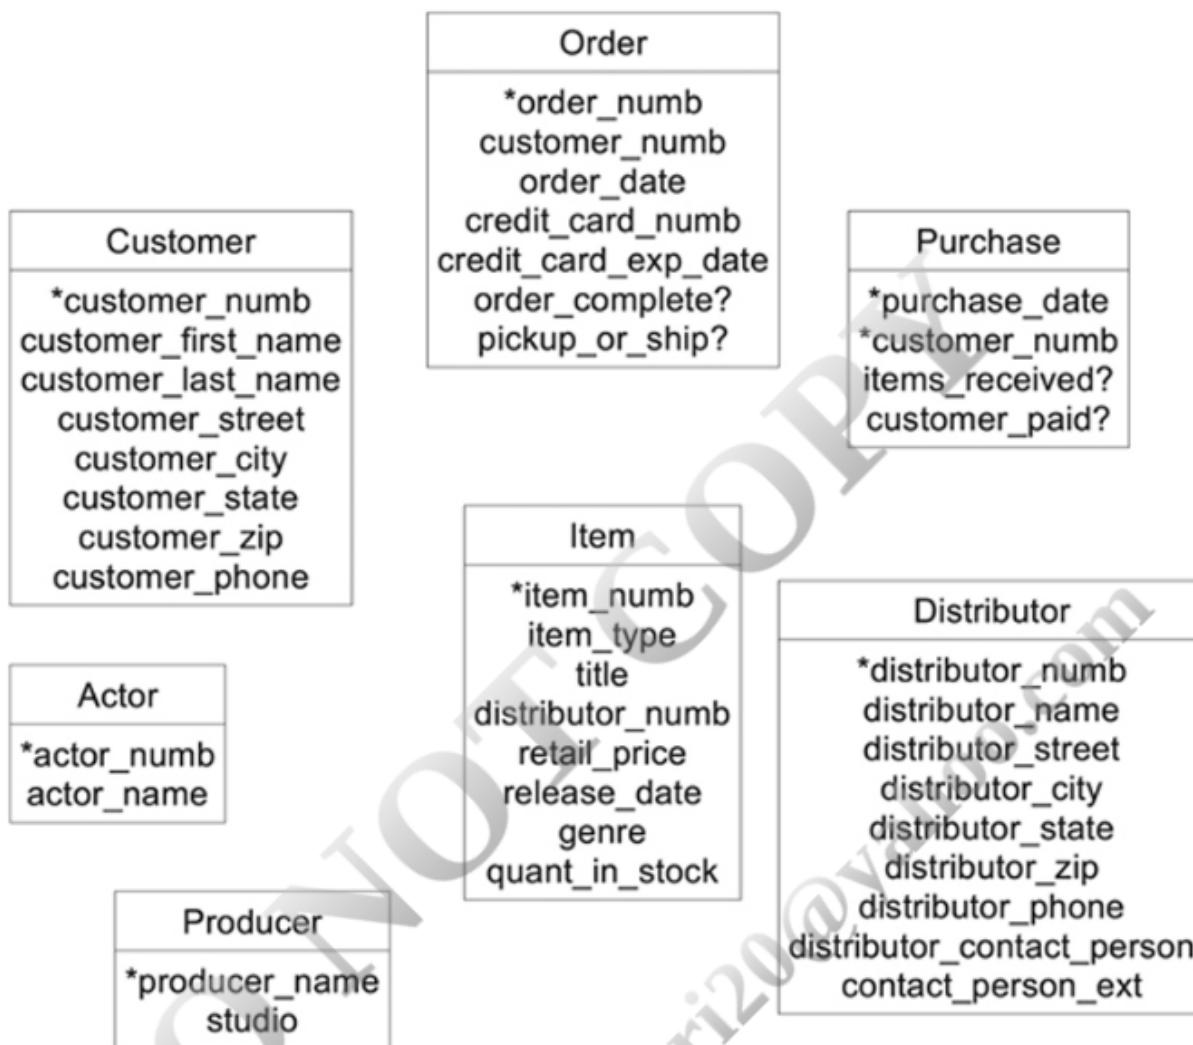


FIGURE 4.4 Major entities and their attributes for the *Antique Opticals* database.

Note: The entities in [Figure 4.4](#) and the remainder of the diagrams in this book were created with a special type of software known as a CASE tool (computer-aided software engineering). CASE tools provide a wide range of data and systems modeling assistance. You will find more details on how CASE tools support the database design process in [Chapter 12](#).

[Figure 4.4](#) demonstrates some of the choices made for the *Antique Opticals* database. Notice, for example, that there is only one entity for merchandise items, yet the store carries new DVDs, new blue-ray discs, and used laserdiscs. The *item_type* attribute distinguishes the three types of merchandise. Because all merchandise items are stored as the same type of entity, queries such as “Show me *Star Wars IV* in any format” will be easy to satisfy using just the item name, and queries such as “Do you have a used *Star Wars IV* laserdisc?” will be easy to satisfy using both the title of the item and its type.

Domains

Each attribute has a *domain*, an expression of the permissible values for that attribute. A domain can be very small. For example, a T-shirt store might have a *Size* attribute for its merchandise items, with the values L, XL, and XXL comprising the entire domain. In contrast, an attribute for a customer’s first name is very large and might be specified only as “text” or “human names.”

A DBMS enforces a domain through a *domain constraint*. Whenever a value is stored in the database, the DBMS verifies that it comes from the attribute’s specified domain. Although in many cases we cannot specify small domains, at the very least the domain assures us that we are getting data of the right type. For example, a DBMS can prevent a user from storing 123 × 50 in an attribute whose domain is currency values. Most DBMSs also provide fairly tight domain checking on date and time attributes, which can help you avoid illegal dates, such as Feb. 30.

Documenting Domains

Printed by: rituadikari20@yahoo.com. Printing is for personal, private use only. No part of this book may be reproduced or transmitted without publisher's prior permission. Violators will be prosecuted.

Documenting Domains

The common formats used for ER diagrams do not usually include domains on the diagrams themselves, but store the domains in an associated document (usually a *data dictionary*, something about which you will read much more throughout this book). However, the version of the Chen method that includes attributes can also include domains underneath each attribute. Notice in [Figure 4.5](#) that three of the domains are fairly general (integer and character), while the domain for the telephone number attribute includes a very specific format. Whether a domain can be constrained in this way depends on the DBMS.

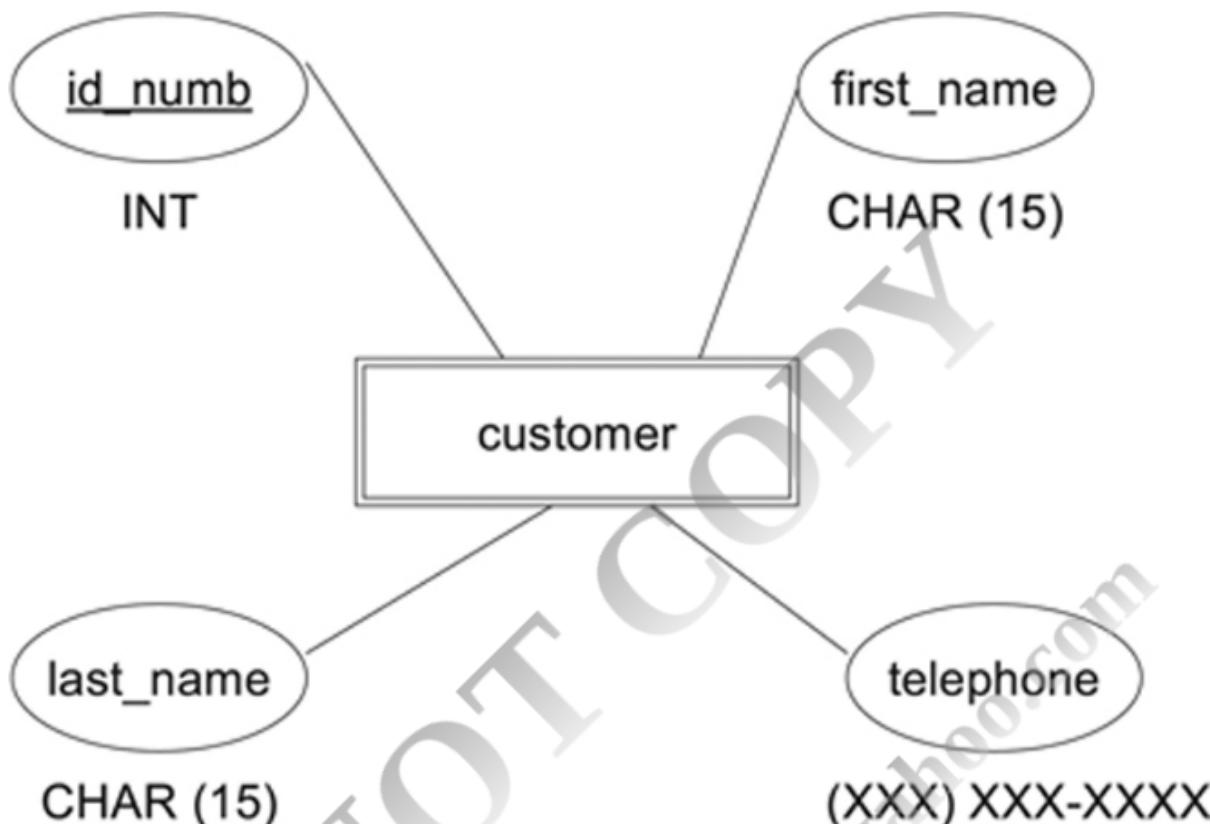


FIGURE 4.5 Indicating domains on an ER diagram.

Note: There is no specific syntax for indicating domains. However, if you know which DBMS you will be using, consider using the data types supported by that product as domains in an ERD to simplify the later conversion to the DBMS's representation.

Practical Domain Choices

The domains that *Antique Opticals* chooses for its attributes should theoretically be independent of the DBMS that the company will use. In practical terms, however, it makes little sense to assign domains that you cannot implement. Therefore, the database designer working for *Antique Opticals* takes a look at the DBMS to see what data types are supported.

Most relational DBMSs that use SQL as their query language provide the following among their data types, any of which can be assigned as a domain to an attribute:

- CHAR: A fixed-length string of text, usually up to 256 characters.
- VARCHAR: A variable-length string of text, usually up to 256 characters.
- INT: An integer, the size of which varies, depending on the operating system.
- DECIMAL and NUMERIC: Real numbers, with fractional portions assigned to the right of the decimal point. When you assign a real number domain, you must specify how many digits the number can contain (including the decimal point) and how many digits should be to the right of the decimal point (the value's *precision*). For example, currency values usually have a precision of two, so a number in the format XXX.XX might have a domain of DECIMAL (6,2).
- DATE: A date.
- TIME: A time.
- DATETIME: The combination of a date and a time.
- BOOLEAN: A logical value (true or false).

Many of today's DBMSs also support a data type known as a binary large object (BLOB) that can store anything binary, such as a graphic.

Choosing the right domain can make a big difference in the accuracy of a database. For example, a US zip code is made up of five or nine digits. Should an attribute for a zip code, therefore, be given a domain of INT? No, for two reasons. First, it would be nice to be able to include the hyphen in nine-digit zip codes. Second, and more importantly, zip codes in the northeast begin with a zero. If they are stored as a number, the leading zero disappears. Therefore, we always choose a CHAR domain for zip codes. Since we never do arithmetic with zip codes, nothing is lost by using character, rather than numeric storage.

By the same token, it is important to choose domains of DATE and TIME for chronological data. As an example, consider what would happen if the dates 01/12/2017 and 08/12/2016 were stored as characters. If you ask the DBMS to choose which date comes first, the DBMS will compare the character strings in alphabetical order and respond the 01/12/2017 comes first, because 01 alphabetically precedes 08. The only way to get character dates to order correctly is to use the format YYYY/MM/DD, a format that is rarely used anywhere in the world. However, if the dates were given a domain of DATE, then the DBMS would order them properly. The DBMS would also be able to perform date arithmetic, finding the interval between two dates or adding constants (for example, the correct number of days in a month) to dates.

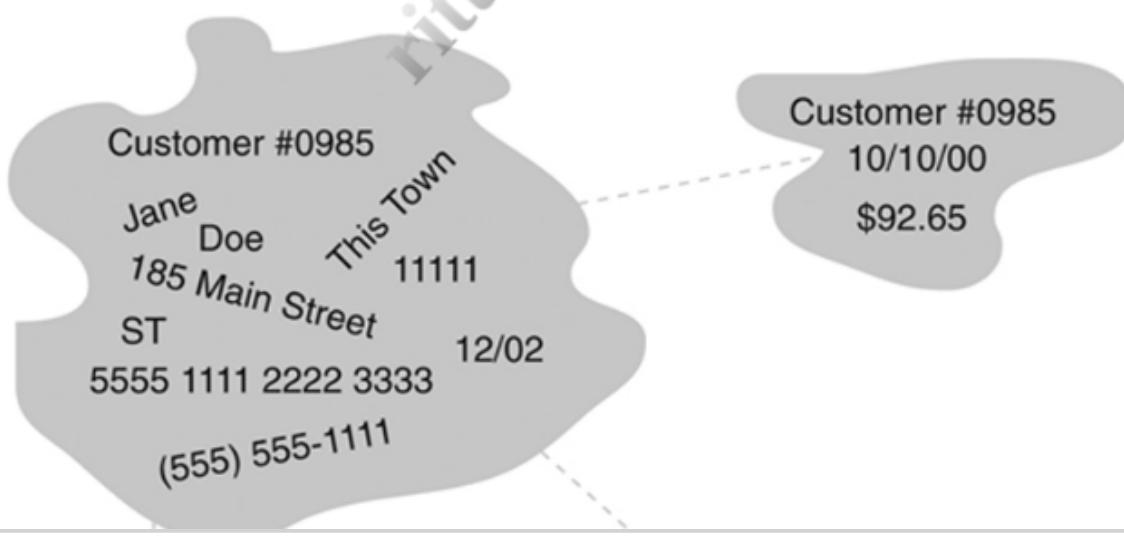
Domains for Numbers

Does it really matter if you choose character or numeric domains for numbers? Yes, it matters a great deal. A computer's internal representation for characters and numbers are very different. When you use an integer or decimal domain, the computer stores a quantity that can be used in computations. Comparisons between values with numeric domains and the ordering of values work in correct numeric order as well. However, if you store a number as a character, the computer encodes the individual digits—typically using the ASCII coding scheme—as text. The ASCII code for a number is not equivalent to the mathematical value of a number. If you try to perform arithmetic using character-coded digits, the result will not be correct. Use character storage for numbers that will never be used in mathematic operations, such as ID numbers and zip codes. If there is even the slightest chance that a computation will be performed using a value, give it a numeric domain.

Basic Data Relationships

Once you have a good idea of the basic entities in your database environment, your next task is to identify the relationships among those entities. There are three basic types of relationships that you may encounter: one-to-one (1:1), one-to-many (1:M), and many-to-many (M:N or M:M).

Before turning to the types of relationships themselves, there is one important thing to keep in mind: The relationships that are stored in a database are between instances of entities. For example, an *Antique Opticals* customer is related to the items he or she orders. Each instance of the customer entity is related to instances of the specific items ordered (Figure 4.6).



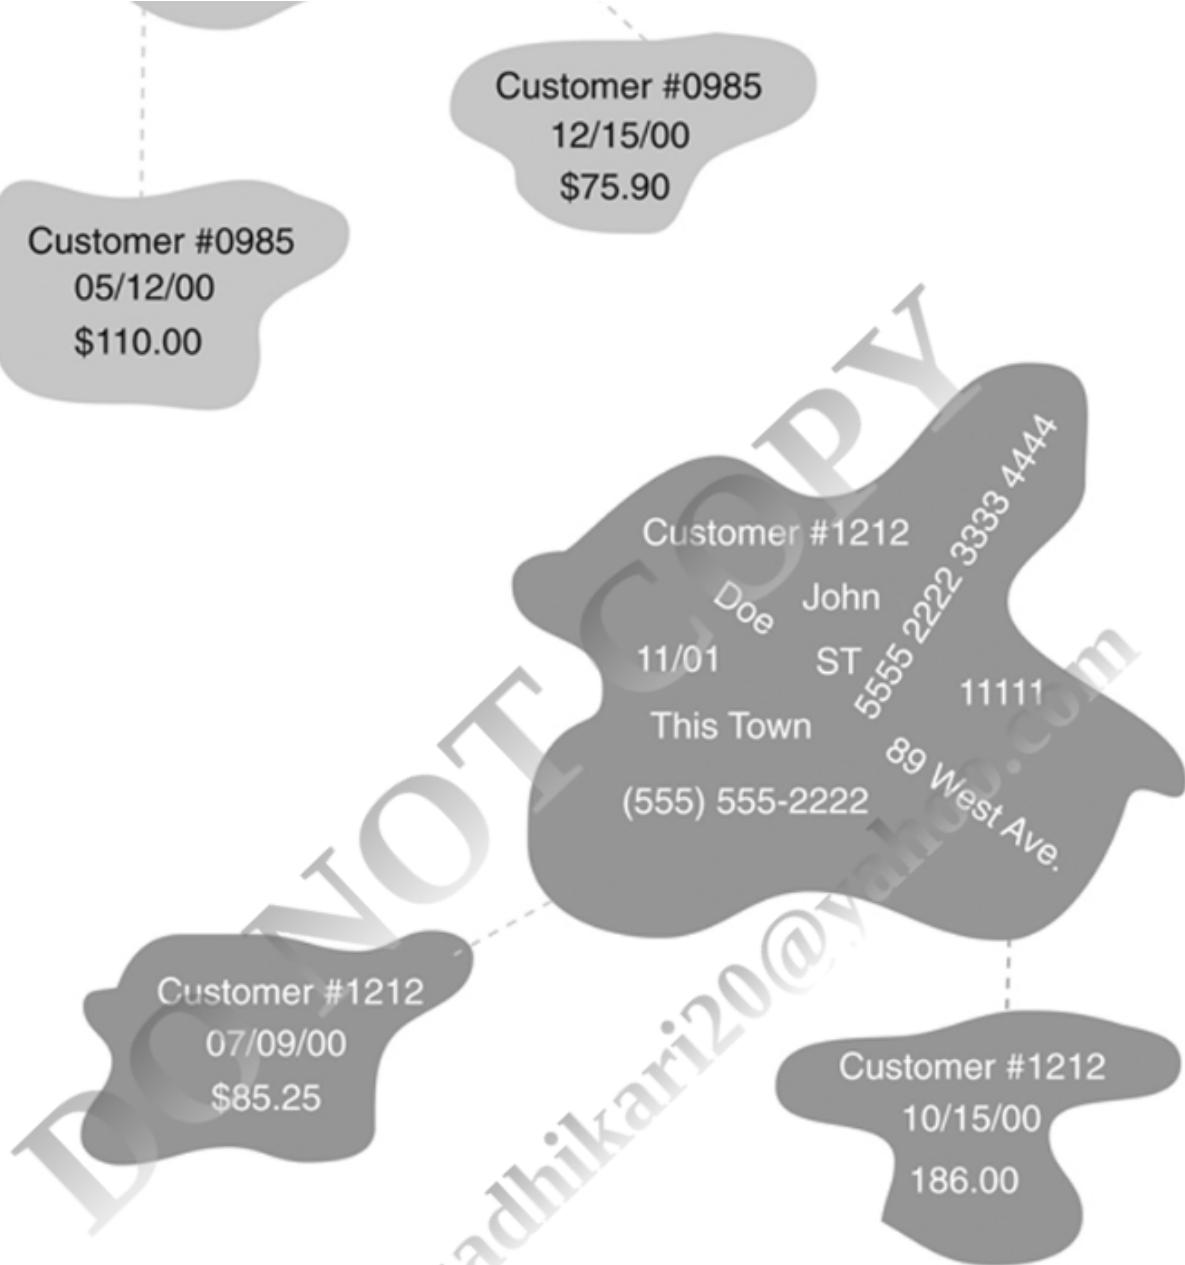


FIGURE 4.6 Relationships between instances of entities in a database.

When we document data relationships, such as when we draw an ER diagram, we show the types of relationships among entities. We are showing the possible relationships that are allowable in the database. Unless we specify that a relationship is mandatory, there is no requirement that every instance of every entity be involved in every documented relationship. For example, *Antique Opticals* could store data about a customer without the customer having any orders to which it is related.

One-to-One Relationships

Consider, for a moment, an airport in a small town and the town in which the airport is located, both of which are described in a database of small town airports. Each of these might be represented by an instance of a different type of entity. The relationship between the two instances can be expressed as “The airport is located in one and only one town, and the town has one and only one airport.”

This is a true *one-to-one relationship* because at no time can a single airport be related to more than one town and no town can be related to more than one airport. (Although there are municipalities that have more than one airport, the towns in the database are too small for that to ever happen.)

If we have instances of two entities (A and B) called A_i and B_j , then a one-to-one relationship exists if at all times A_i is related to no instances of entity B or one instance of entity B, and B_j is related to no instances of entity A or one instance of entity A.

True one-to-one relationships are very rare in business. For example, assume that *Antique Opticals* decides to start dealing with a new distributor of DVDs. At first, the company orders only one specialty title from the new distributor. If we peered inside the database, we would see that the instance of the distributor entity was related to just the one merchandise item instance. This would then appear to be a one-to-one relationship. However, over time, *Antique Opticals* may choose to order more titles from the new distributor, which would violate the rule that the distributor must be related to no more than one merchandise item. (This is an example of a one-to-many relationship, which is discussed in the next section of this chapter.)

By the same token, what if *Antique Opticals* created a special credit card entity to hold data about the credit cards that customers use to place

Printed by: rituadhibkari20@yahoo.com. Printing is for personal, private use only. No part of this book may be reproduced or transmitted without publisher's prior permission. Violators will be prosecuted.

By the same token, what if *Antique Opticals* created a special credit card entity to hold data about the credit cards that customers use to place orders? Each order can be charged to only one credit card. There would therefore seem to be a one-to-one relationship between an instance of an order entity and an instance of a credit card entity. However, in this case, we are really dealing with a single entity. The credit card number and the credit card's expiration date can become attributes of the order entity. Given that only one credit card is allowed per order, the attributes are not multivalued; no separate entity is needed.

If you think you are dealing with a one-to-one relationship, look at it very carefully. Be sure that you are not really dealing with a special case of a one-to-many relationship or two entities that should really be one.

One-to-Many Relationships

The most common type of relationship is a *one-to-many relationship*. (In fact, relational databases are constructed from the rare one-to-one relationship and numerous one-to-many relationships.) For example, *Antique Opticals* typically orders many titles from each distributor, and a given title comes from only one distributor. By the same token, a customer places many orders, but an order comes from only one customer.

If we have instances of two entities (A and B), then a one-to-many relationship exists between two instances (A_i and B_j) if A_i is related to zero, one, or more instances of entity B, and B_j is related to zero or one instance of entity A.

Other one-to-many relationships include that between a daughter and her biological mother. A woman may have zero, one, or more biological daughters; a daughter can have only one biological mother. As another example, consider a computer and its CPU. A CPU may not be installed in any computer, or it may be installed in at most one computer; a computer may have no CPU, one CPU, or more than one CPU.

The example about which you read earlier, concerning *Antique Opticals* and the distributor from which the company ordered only one title, is actually a one-to-many relationship where the "many" is currently "one." Remember that, when we are specifying data relationships, we are indicating possible relationships, and not necessarily requiring that all instances of all entities participate in every documented relationship. There is absolutely no requirement that a distributor be related to any merchandise item, much less one or more merchandise items. (It might not make much sense to have a distributor in the database from which the company does not order, but there is nothing to prevent data about that distributor from being stored.)

Many-to-Many Relationships

Many-to-many relationships are also very common. There is, for example, a many-to-many relationship between an order placed by an *Antique Opticals* customer and the merchandise items carried by the store. An order can contain multiple items; each item can appear on more than one order. The same is true of the orders placed with distributors. An order can contain multiple items, and each item can appear on more than one order.

A many-to-many relationship exists between entities A and B if for two instances of those entities (A_i and B_j), A_i can be related to zero, one, or more instances of entity B, and B_j can be related to zero, one, or more instances of entity A.

Many-to-many relationships present two major problems to a database's design. These issues, and the way in which we solve them, are discussed in the next major section of this chapter (Dealing with Many-to-Many Relationships).

Weak Entities and Mandatory Relationships

As we have been discussing types of data relationships, we have defined those relationships by starting each with "zero," indicating that the participation by a given instance of an entity in a relationship is optional. For example, *Antique Opticals* can store data about a customer in its database before the customer places an order. Therefore, an instance of the customer entity does not have to be related to any instances of the order entity.

However, the reverse is not true in this database. An order *must* be related to a customer. Without a customer, an order cannot exist. An order is, therefore, an example of a *weak entity*, one that cannot exist in the database, unless a related instance of another entity is present and related to it. An instance of the customer entity can be related to zero, one, or more orders. However, an instance of the order entity must be related to one and only one customer. The "zero" option is not available to a weak entity. The relationship between an instance of the order entity and an instance of the customer entity is, therefore, a *mandatory relationship*.

Identifying weak entities and their associated mandatory relationships can be very important for maintaining the consistency and integrity of the database. Consider the effect, for example, of storing an order without knowing the customer to whom it belongs. There would be no way to ship the item to the customer, causing a company to lose business.

By the same token, we typically specify the relationship between an order and the order lines (the specific items on the order) as mandatory, because we don't want to allow an order line to exist in the database without it being related to an order. (An order line is meaningless without knowing the order to which it belongs.)

In contrast, we can allow a merchandise item to exist in a database without indicating the supplier from which it comes (assuming that there is only one source per item). This lets us store data about new items before we have decided on a supplier. In this case, the relationship between a supplier and an item is not mandatory (sometimes described as zero-to-many, rather than one-to-many).

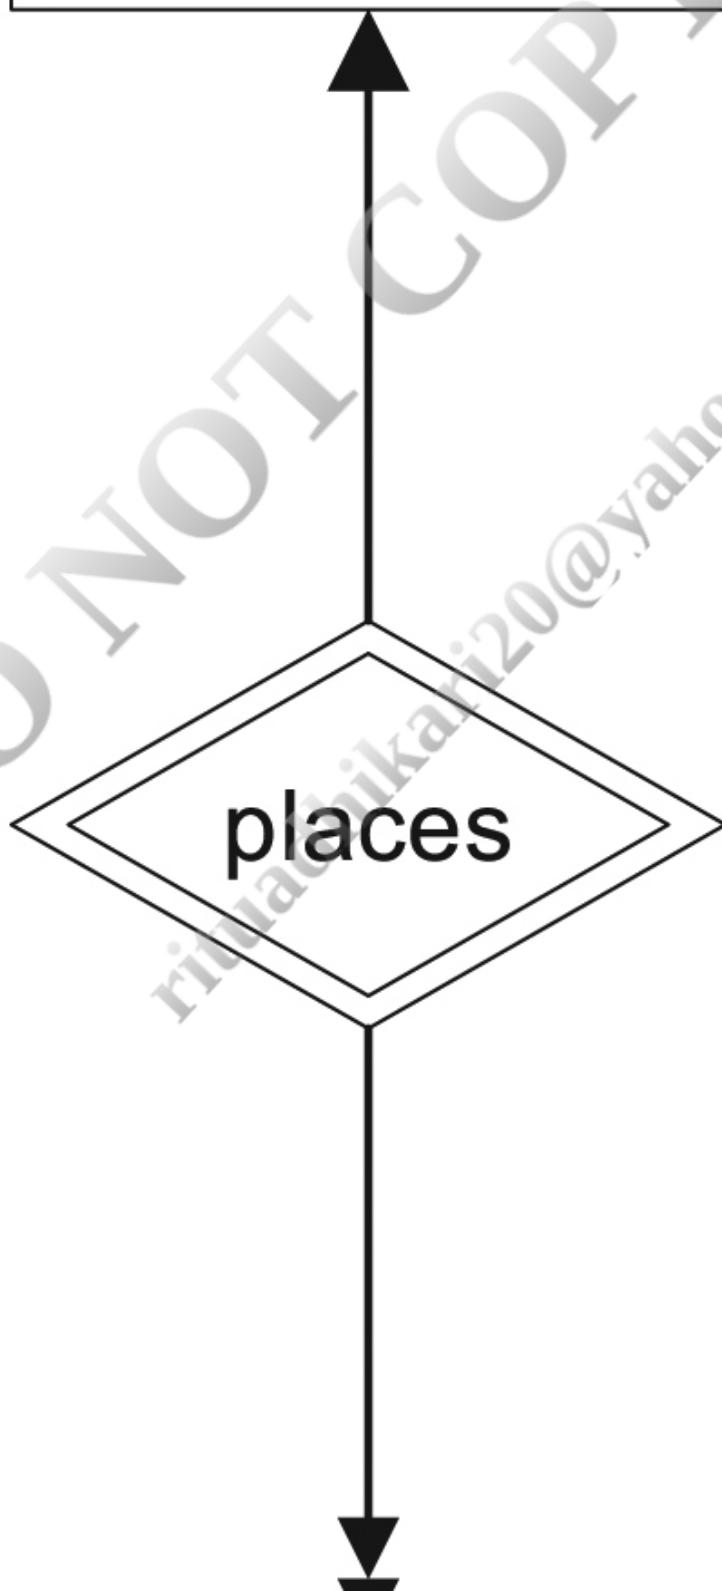
Documenting Relationships

The Chen, IE, and UML methods of drawing ER diagrams have different ways of representing relationships, each of which has its advantages in terms of the amount of information it provides and its complexity.

The Chen Method

The Chen method uses diamonds for relationships and lines with arrows to show the types of relationships between entities. For example, in Figure 4.7 you can see the relationship between an *Antique Opticals* customer and an order. The single arrow pointing toward the customer entity indicates that an order belongs to at most one customer. The double arrow pointing toward the order entity indicates that a customer can place one or more orders. The word within the relationship diamond gives some indication of the meaning of the relationship.

Customer



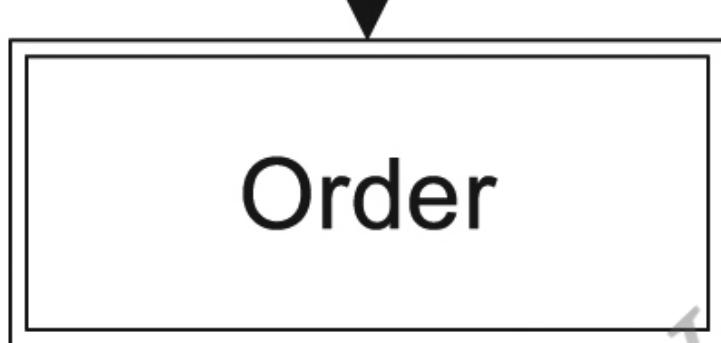
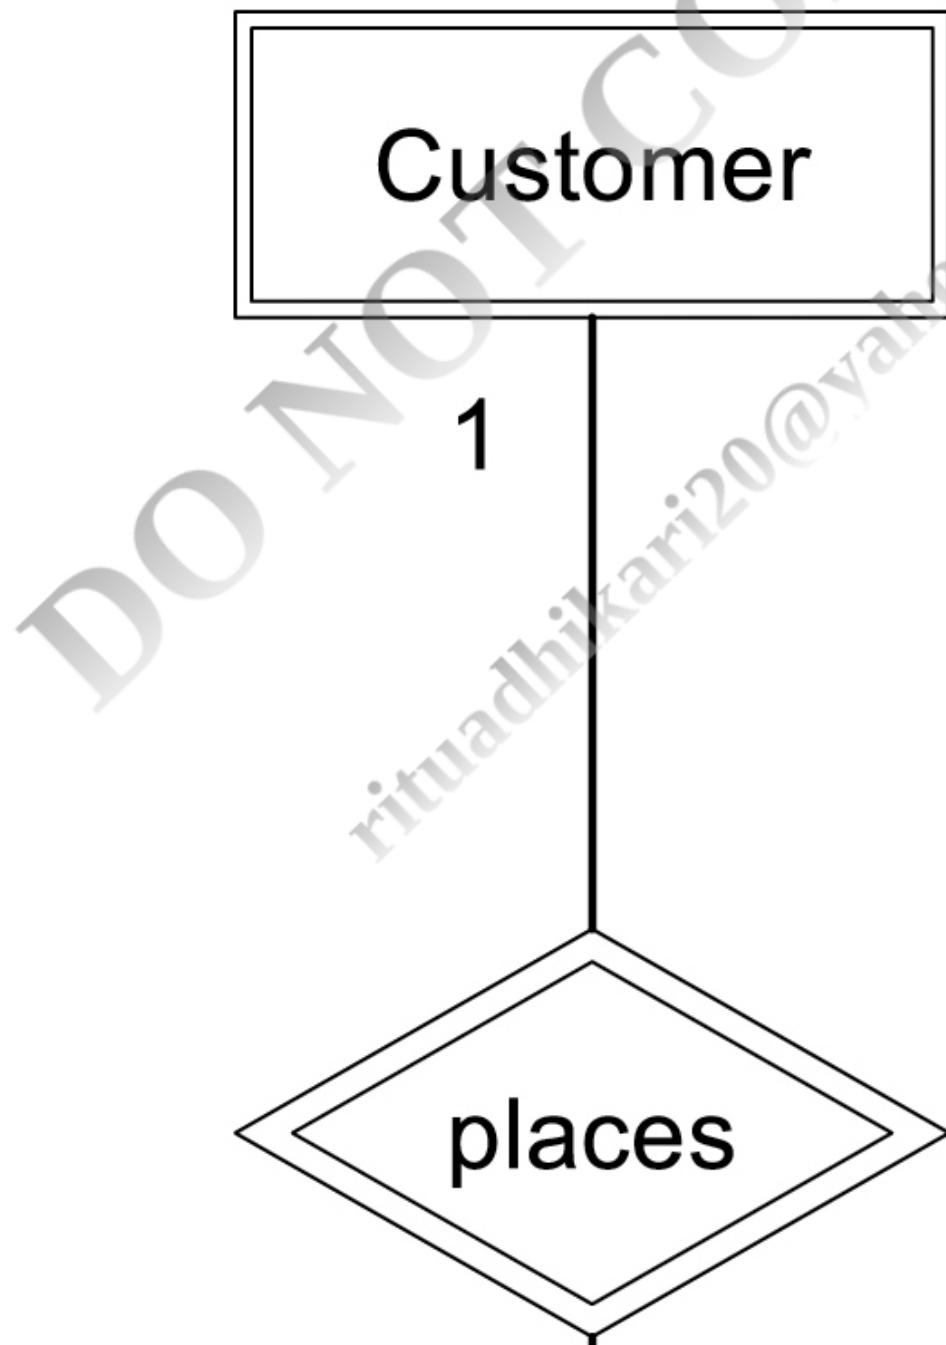


FIGURE 4.7 Using the Chen method with relationship diamonds and arrows.

There are two alternative styles within the Chen method. The first (for example, [Figure 4.8](#)) replaces the arrows with numbers and letters. A "1" indicates that an order comes from one customer. The "M" (or an "N") indicates that a customer can place many orders.



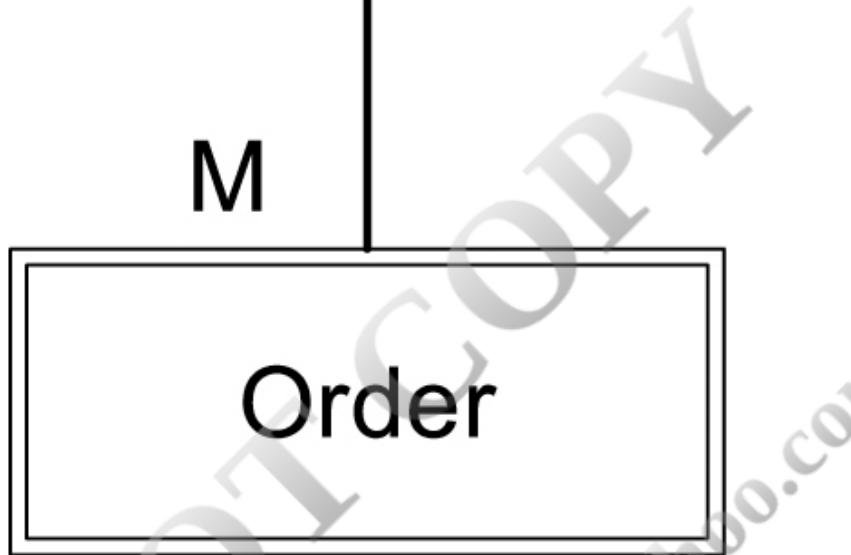
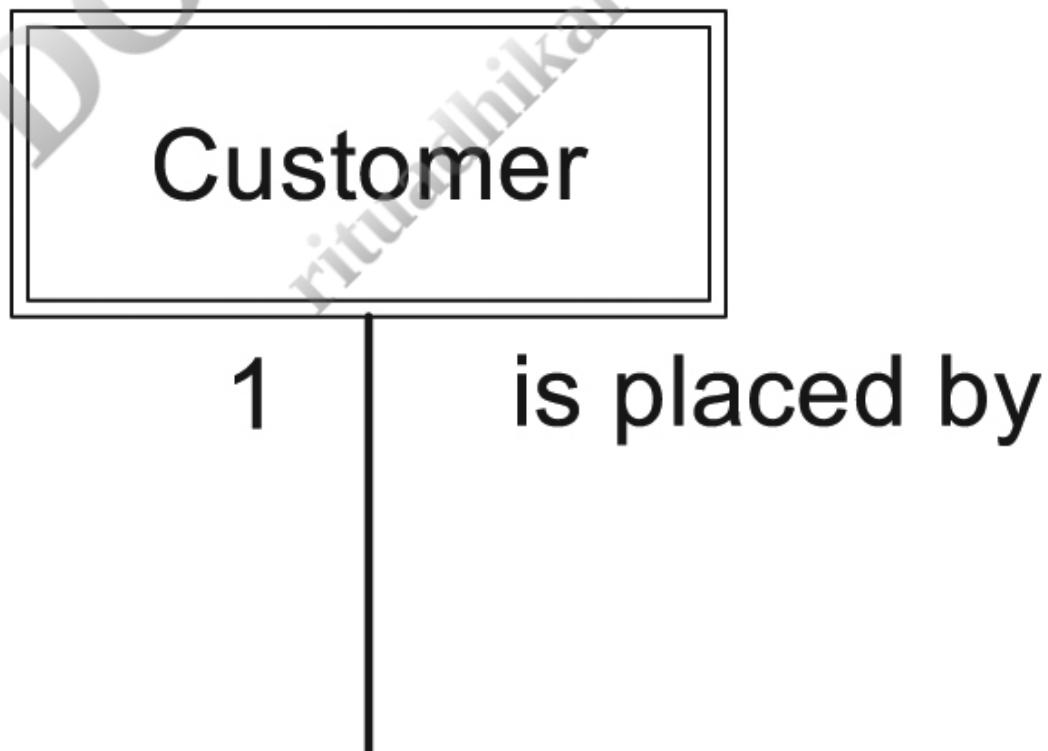


FIGURE 4.8 A Chen ER diagram using letters and numbers rather than arrows to show relationships.

The second alternative addresses the problem of trying to read the relationship in both directions when the name of the relationship is within the diamond. "Customer places order" makes sense, but "order places customer" does not. To solve the problem, this alternative removes the relationship name from the diamond and adds both the relationship and its inverse to the diagram, as in Figure 4.9. This version of the diagram can be read easily in either direction: "A customer places many orders" and "An order is placed by one customer."



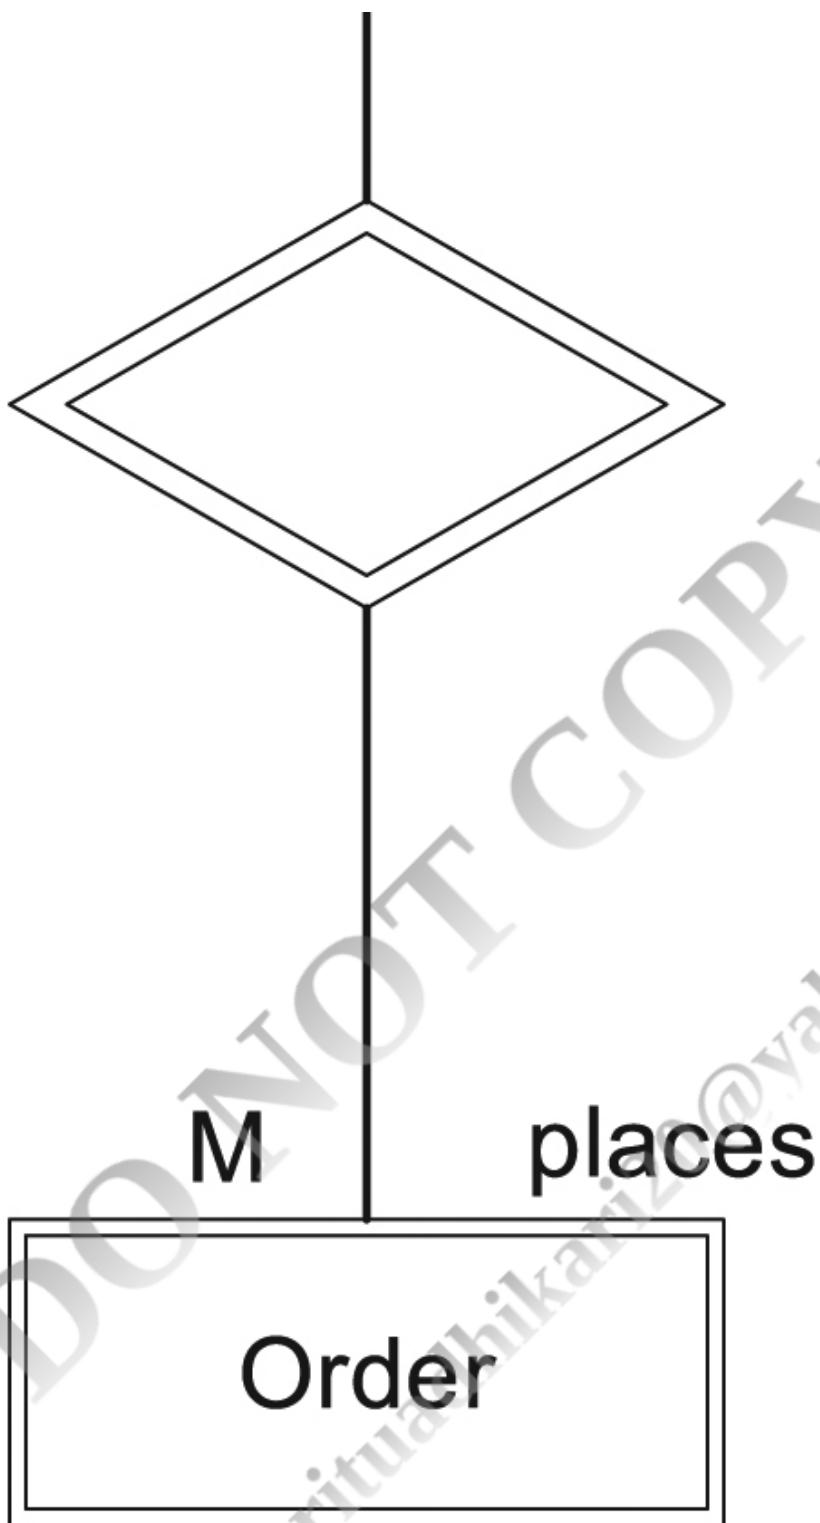


FIGURE 4.9 Adding inverse relationships to a Chen method ER diagram.

There is one major limitation to the Chen method of drawing ER diagrams: There is no obvious way to indicate weak entities and mandatory relationships. For example, an order should not exist in the database without a customer. Therefore, order is a weak entity, and its relationship with a customer is mandatory.

IE Style Diagrams

The IE diagramming style exchanges simplicity in line ends for added information. As a first example, consider Figure 4.10. This is the same one-to-many relationship we have been using to demonstrate the Chen method ER diagrams. However, in this case, the ends of the lines (some of which look a bit like “crow’s feet”) indicate which relationships are mandatory.

Customer

*customer_numb
customer_first_name
customer_last_name
customer_street
customer_city
customer_state
customer_zip
customer_phone

+



Order

***order_num
customer_num
order_date
credit_card_num
credit_card_exp_date
order_complete?
pickup_or_ship?**

FIGURE 4.10 A one-to-many relationship using the IE method.

The double line below the customer entity means that each order is related to one and only one customer. Because zero is not an option, the relationship is mandatory. In contrast, the 0 and crow's foot connected to the order entity mean that a customer may have zero, one, or more orders.

There are four symbols used at the ends of lines in an IE diagram:

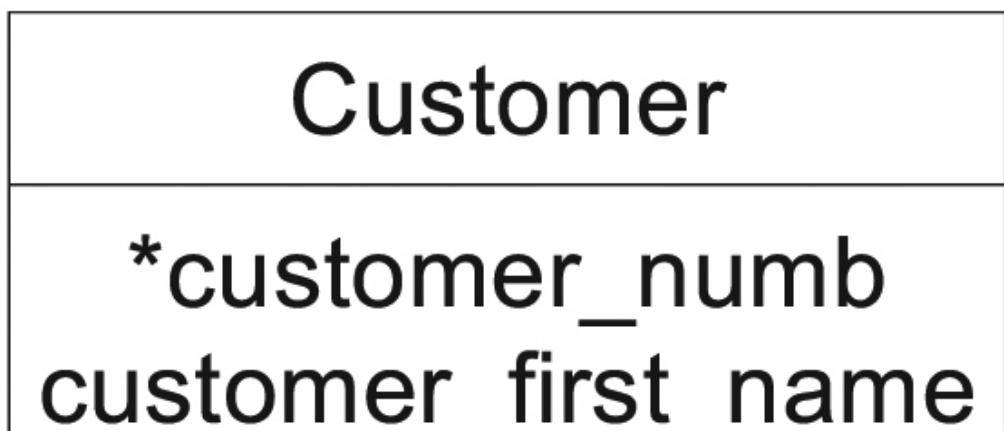
- | : One and one only (mandatory relationship)
- 0| : Zero or one
- >1: One or more (mandatory relationship)
- >0: Zero, one, or more

Although we often see the symbols turned 90 degrees, as they are in [Figure 4.10](#), they are actually readable if viewed sideways as in the preceding list.

An IE method ER diagram often includes attributes directly on the diagram. As you can see in [Figure 4.10](#), entity identifiers are marked with an asterisk.

UML Style Diagrams

UML notation for entity relationships is very similar to IE notation. However, the symbols at the ends of lines are replaced by numeric representations of the type of relationship ([Figure 4.11](#)).



customer_last_name
customer_street
customer_city
customer_state
customer_zip
customer_phone

1

0...*

Order

***order_num**
customer_num
order_date

`credit_card_num`
`credit_card_exp_date`
`order_complete?`
`pickup_or_ship?`

FIGURE 4.11 A one-to-many relationship using UML notation.

The possible relationships are:

- 1: One and only one (mandatory)
- 1...*: One or more (mandatory)
- 0...1: Zero or one
- 0...*: Zero, one, or more

Basic Relationships for Antique Opticals

The major entities in the *Antique Opticals* database are diagrammed in Figure 4.12. You read the relationships in the following way:

- One customer can place zero, one or more orders. An order comes from one and only one customer.
- The store may make many purchases of used discs from each customer. A purchase transaction comes from one and only one customer.
- An order has one or more items on it. An item can appear in zero, one, or more orders.
- A purchase is made up of one or more items. An item can be purchased zero, one, or more times.
- An actor appears in zero, one, or more items. An item has zero, one, or more actors in it. (There may occasionally be films that feature animals rather than human actors; therefore, it is probably unwise to require that every merchandise item be related to at least one actor.)
- Each item has zero, one, or more producers. Each producer is responsible for zero, one, or more items. (Although in practice you would not store data about a producer unless that producer was related to an item, leaving the relationship between a producer and an item as optional means that you can store producers without items, if necessary.)
- Each item comes from zero or one distributor. Each distributor supplies zero, one, or more items.

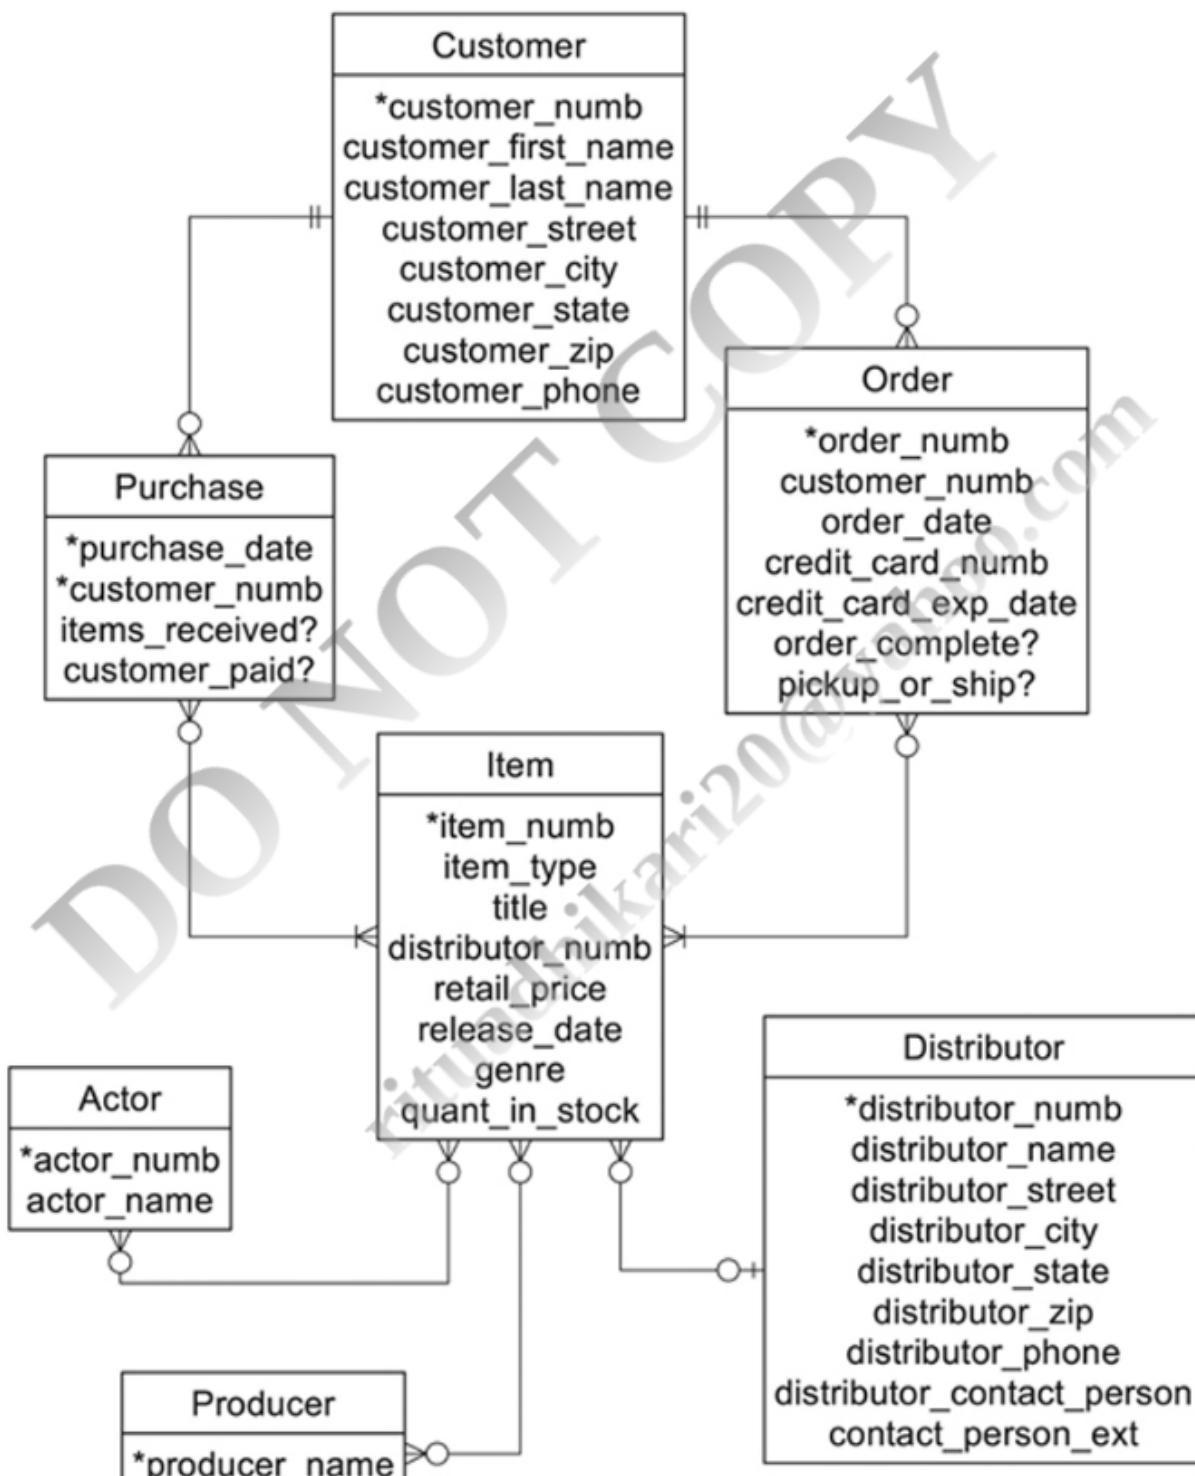


FIGURE 4.12 The major entities and the relationships between them in the *Antique Opticals* database.

The major thing to notice about this design is that there are four many-to-many relationships: order to item, purchase to item, actor to item, and producer to item. Before you can map this data model to a relational database, they must be handled in some way.

Dealing with Many-to-Many Relationships

As you read earlier, there are problems with many-to-many relationships. The first is fairly straightforward: The relational data model cannot handle many-to-many relationships directly; it is limited to one-to-one and one-to-many relationships. This means that you must replace the many-to-many relationships that you have identified in your database environment with a collection of one-to-many relationships if you want to be able to use a relational DBMS.

The second is a bit more subtle. To understand it, consider the relationship between an order an *Antique Opticals* customer places with the store and the merchandise items on the order. There is a many-to-many relationship between the order and the item because each order can be for many items, and each item can appear on many orders (typically orders from different customers). Whenever a customer orders an item, the number of copies of the item varies, depending on how many copies the customer needs. (Yes, typically people order only one copy of a movie, but we need to allow them to order as many as they want. It certainly would be poor business practice to do otherwise.)

Now the question: Where should we store the quantity being ordered? It cannot be part of the order entity, because the quantity depends on the item being ordered. By the same token, the quantity cannot be part of the item entity, because the quantity depends on the specific order.

What you have just seen is known as *relationship data*, data that apply to the relationship between two entities, rather than to the entities themselves. In the relational data model, however, relationships cannot have attributes. We must therefore have some entity to represent the relationship between the two entities participating in the M:M relationship, an entity to which the relationship data can belong.

Composite Entities

Entities that exist to represent the relationship between two or more other entities are known as *composite entities*. As an example of how composite entities work, consider once again the relationship between an order placed by an *Antique Opticals* customer and the items on that order.

What we need is an entity that tells us that a specific title appears on a specific order. If you look at Figure 4.13, you will see three order instances, and three merchandise item instances. The first order for customer 0985 (Order #1) contains only one item (item 09244). The second order for customer 0985 (Order #2) contains a second copy of item 09244, as well as item 10101. Order #3, which belongs to customer 1212, also has two items on it (item 10101 and item 00250).

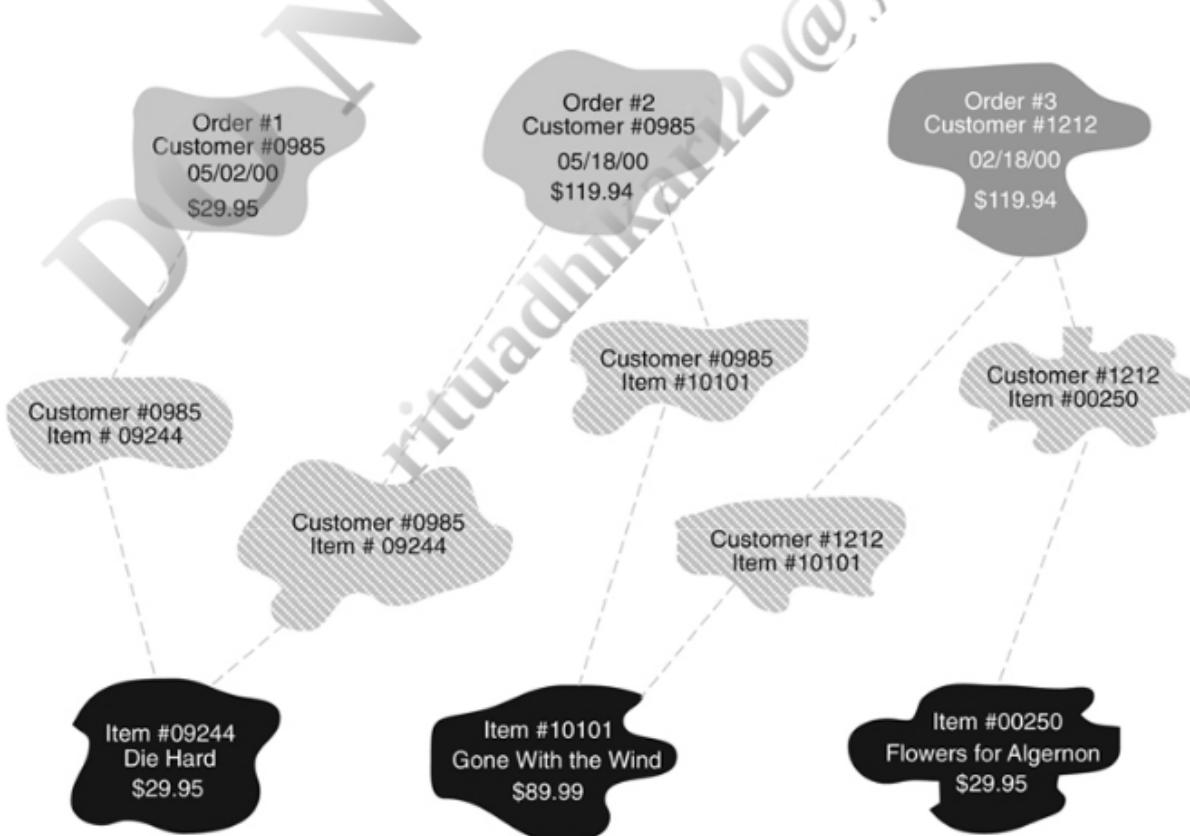


FIGURE 4.13 Using instances of composite entities to change many-to-many relationships into one-to-many relationships.

There are five items ordered among the three orders. The middle of the diagram, therefore, contains five instances of a composite entity we will

Printed by: rituadhidhikari20@yahoo.com. Printing is for personal, private use only. No part of this book may be reproduced or transmitted without publisher's prior permission. Violators will be prosecuted.

call a “line item” (thinking of it as a line item on a packing slip or invoice). The line item entity has been created solely to represent the relationship between an order and a merchandise item.

Each order is related to one line item instance for each item on the order. In turn, each item is related to one line item instance for each order on which it appears. Each line item instance is related to one and only one order; it is also related to one and only one merchandise item. As a result, the relationship between an order and its line items is one-to-many (one order has many line items) and the relationship between an item and the orders on which it appears is one-to-many (one merchandise item appears in many line items). The presence of the composite entity has removed the original many-to-many relationship.

If necessary, the composite entity can be used to store relationship data. In the preceding example, we might include an attribute for the quantity ordered, a flag to indicate whether it has been shipped, and a shipping date.

Documenting Composite Entities

In some extensions of the Chen method for drawing ER diagrams, the symbol for a composite entity is the combination of the rectangle used for an entity, and the diamond used for a relationship:

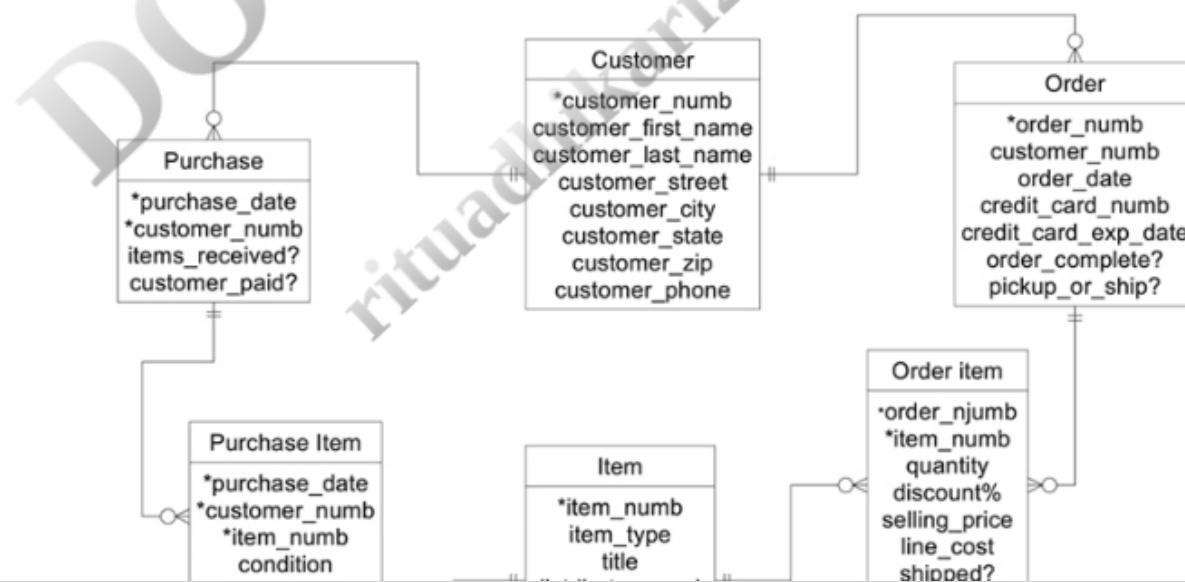


The IE and UML styles, however, have no special symbol for a composite entity.

Resolving *Antique Opticals*' Many-to-Many Relationships

To eliminate *Antique Opticals*' many-to-many relationships, the database designer must replace each many-to-many relationship with a composite entity, and two one-to-many relationships. As you can see in Figure 4.14, the four new entities are as follows:

- Order item: An instance of the order item entity represents one item appearing on one order. Each order can have many “order items,” but an ordered item must appear on one and only one order. By the same token, an ordered item contains one and only one item, but the same item can appear in many order item instances, each of which is related to a different order.
- Purchase item: An instance of the purchase item entity represents one used laser disc purchased from one customer as part of a purchase of one or more discs. Many items can be purchased during a single transaction, but each item purchased is purchased during only one transaction. The purpose of the purchase item entity is therefore the same as the order item entity: to represent specific items in a single transaction.
- Performance: The performance entity represents one actor appearing in one film. Each performance is for one and only one film, although a film can have many performances (one for each actor in the film). Conversely, an actor is related to one performance for each film in which he or she appears, although each performance is in one and only one film.
- Production: The production entity represents one producer working on one film. A producer may be involved in many productions, although each production relates to one and only one producer. The relationship with item indicates that each film can be produced by many producers, but that each production relates to only one item.



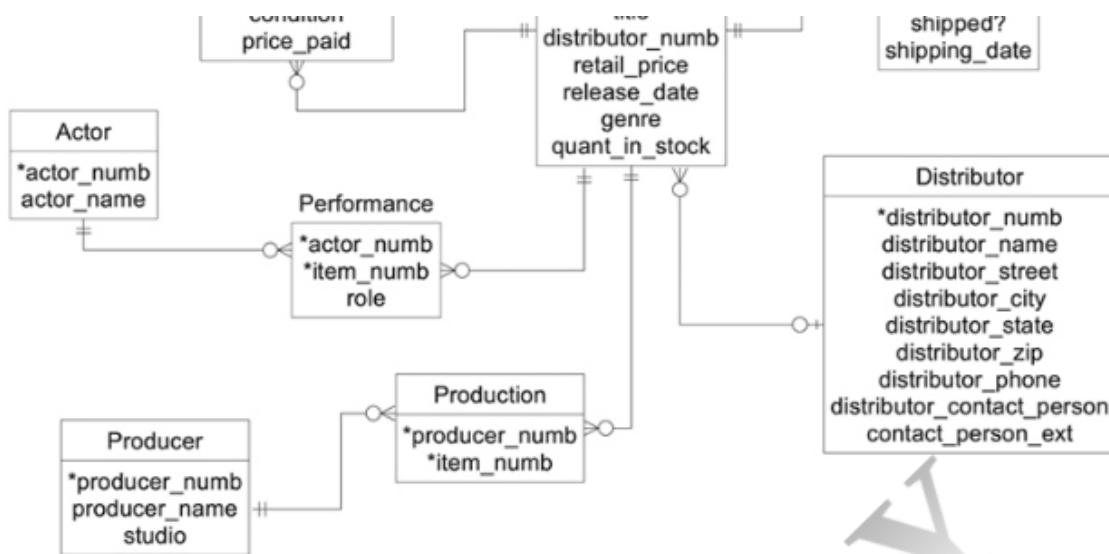


FIGURE 4.14 The complete ER diagram for the *Antique Opticals* database.

Note: If you find sorting out the relationships in Figure 4.14 a bit difficult, keep in mind that if you rotate the up-and-down symbols 90 degrees, you will actually be able to read the relationships from the symbols at the ends of the lines.

Because composite entities exist primarily to indicate a relationship between two other entities, they must be related to both of their parent entities. This is why the relationship between each composite entity in Figure 4.14 and its parents is mandatory.

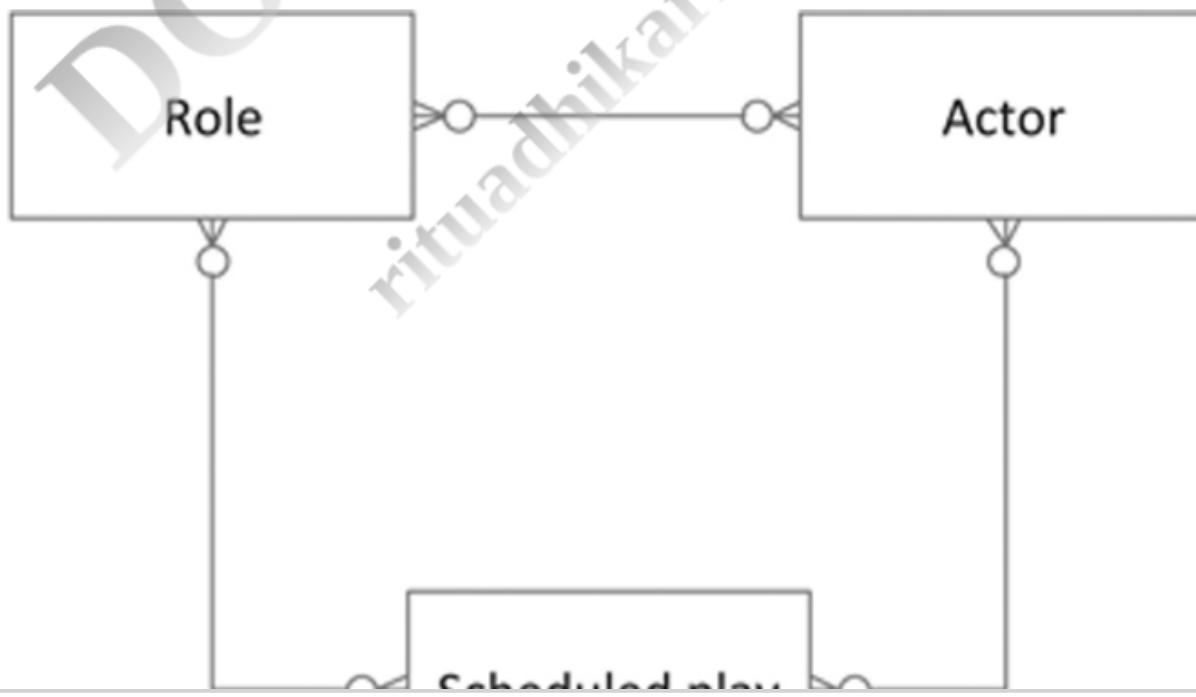
N-Way Composite Entities

There is no reason a composite entity can't represent a relationship between more than two parent entities. As an example, assume that we are creating a database for a community theater group. We want to keep track of which actors performed which roles during the run of each play. To do this, we need entities for the roles, the actors, and the scheduled plays. A scheduled play is a specific play performed at a specific location, during a specific range of dates. (There are other entities in the database, such as the theater and the play, but they don't factor in to this discussion.)

The relationships are as follows:

- A role may be performed by many actors, and an actor may perform many roles.
- An actor may appear in many scheduled plays, and a scheduled play may have many actors.
- A role may appear in many scheduled plays, and a scheduled play may have many roles.

When we draw the ER diagram, we end up with circular many-to-many relationships (Figure 4.15). These relationships are indeed correct, but they don't capture one very important piece of information: which actor performed which role, in which scheduled play. The M:M relationship between a role and an actor tells us who performed which role, but can't tell us the scheduled play. The M:M relationship between an actor and a scheduled play tells us who was in each performance, but can't tell us the role the actor portrayed. Finally, the M:M relationship between a role and a scheduled play can tell us what the roles were in a scheduled play, but not who performed each role during a specific performance.



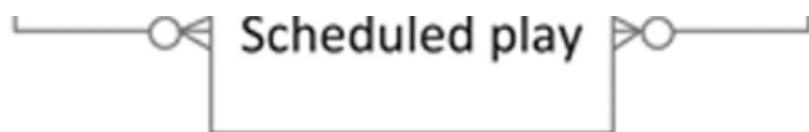


FIGURE 4.15 Circular many-to-many relationships.

A performance—who performed what role, during what scheduled play—needs three pieces of data to describe it: the actor, the role, and the scheduled play. The only way to capture this is to create a single composite entity that has for its parents one actor, one role, and one scheduled play. As you can see in [Figure 4.16](#), when we introduce the performance entity, all the relationships become 1:M, with the performance on the “many” end. Each performance, therefore, represents one actor performing one role in one scheduled play.

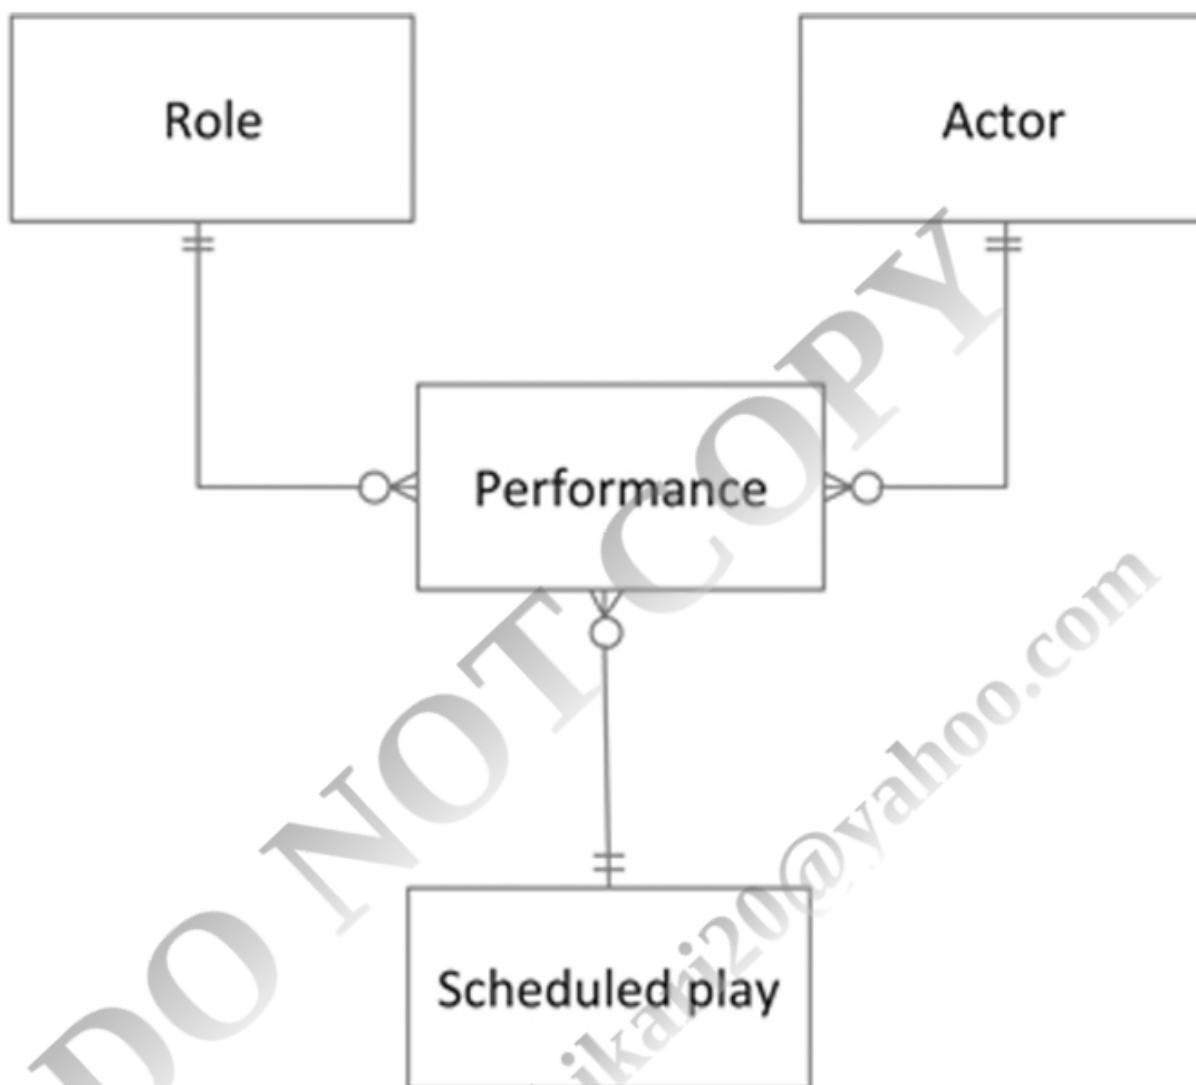


FIGURE 4.16 A three-way composite entity.

Relationships and Business Rules

In many ways, database design is as much an art as a science. Exactly what the “correct” design for a specific business happens to be depends on the business rules; what is correct for one organization may not be correct for another.

As an example, assume that you are creating a database for a small establishment that has more than one store. One of the things you are being asked to model in the database is an employee’s schedule. Before you can do that, you need to answer the question of the relationship between an employee and a store: Is it one-to-many or many-to-many? Does an employee always work at one store—in which case the relationship is one-to-many—or can an employee split his or her time between more than one store, producing a many-to-many relationship? You can see two possible relationships in [Figure 4.17](#). The diagram on the left says that a store has one or more employees, and an employee can work at zero or one store. In contrast, the diagram on the right says that a store may have zero or more employees, and that an employee may work at zero or more stores. Choosing one is not a matter of right or wrong database design, but an issue of how the business operates.

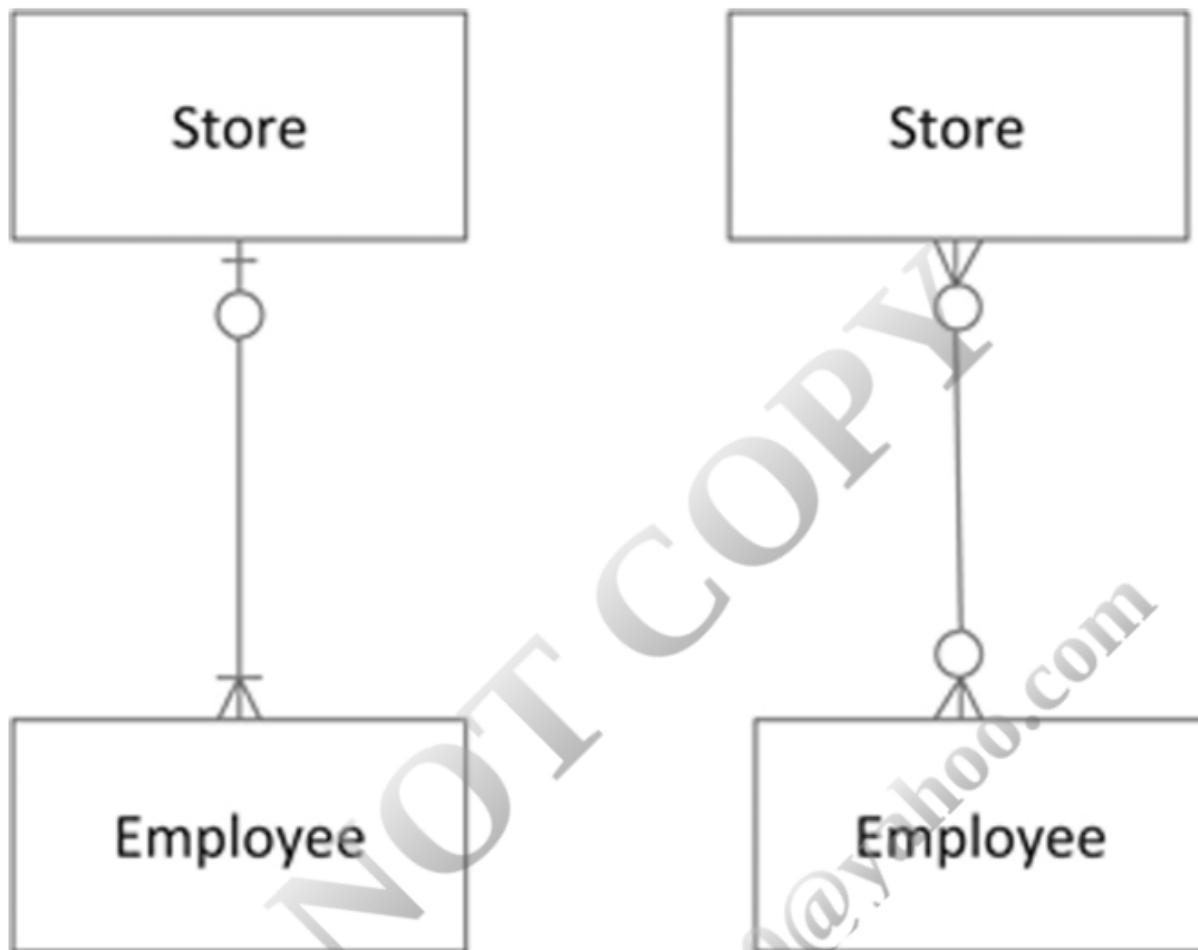
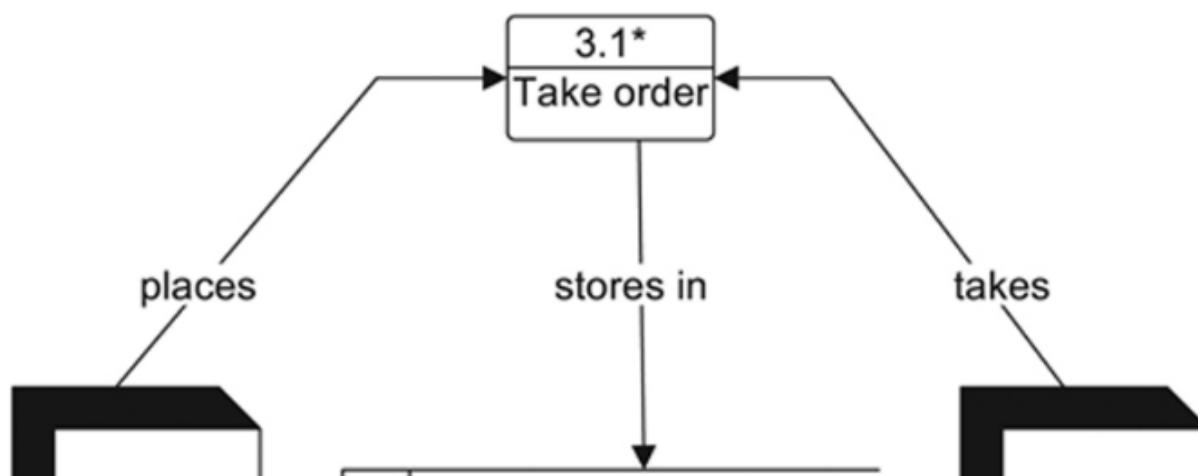


FIGURE 4.17 Two possible relationships requiring a decision based on business rules.

Data Modeling Versus Data Flow

One of the most common mistakes people make when they are beginning to do data modeling is to confuse data models with data flows. A *data flow* shows how data are handled within an organization, including who handles the data, where the data are stored, and what is done to the data. In contrast, a *data model* depicts the internal, logical relationships between the data, without regard to who is handling the data or what is being done with them.

Data flows are often documented in *data flow diagrams* (DFDs). For example, in [Figure 4.18](#) you can see a top-level data flow diagram for *Antique Opticals*. The squares with drop shadows represent the people who are handling the data. Simple rectangles with numbers in them represent processes, or things that are done with the data. A place where data are stored (a *data store*) appears as two parallel lines, in this example containing the words "Main database." The arrows on the lines in the diagram show the way in which data pass from one place to another. For example, "Take order" is a process. Two people generate the data, the Customer and the Employee. Notice that the arrows on the lines between those two and the "Take order" process point to the process, showing that the data flows from the people to the process.



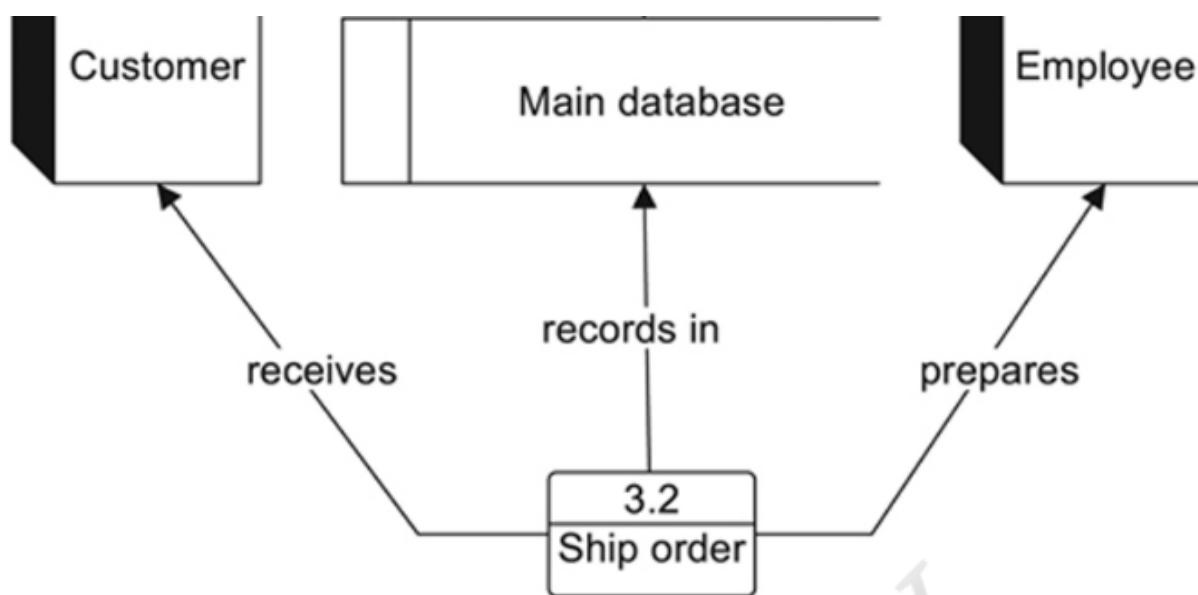
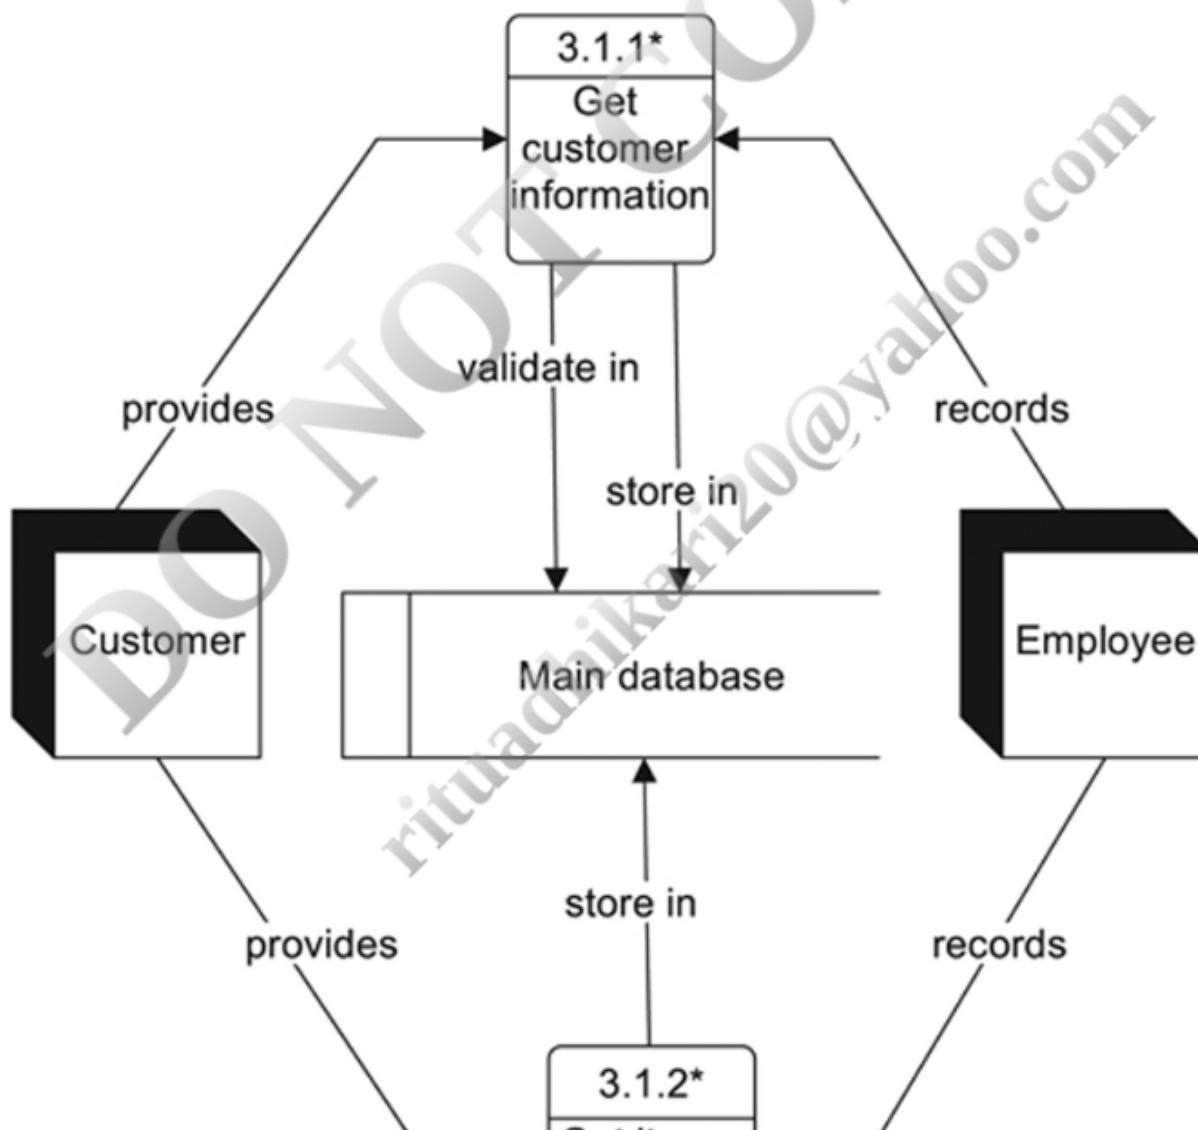


FIGURE 4.18 A top-level data flow diagram for *Antique Opticals*.

Data flow diagrams are often exploded to show more detail. For example, Figure 4.19 contains an explosion of the “Take order” process from Figure 4.18. You can now see that the process of taking an order involves two major steps: getting customer information, and getting item information.



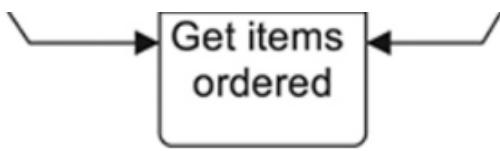


FIGURE 4.19 An explosion of the “Take order” process from [Figure 4.18](#).

Each of the processes in [Figure 4.19](#) can be exploded even further to show additional detail ([Figures 4.20 and 4.21](#)). At this point, the diagrams are almost detailed enough so that an application designer can plan an application program.

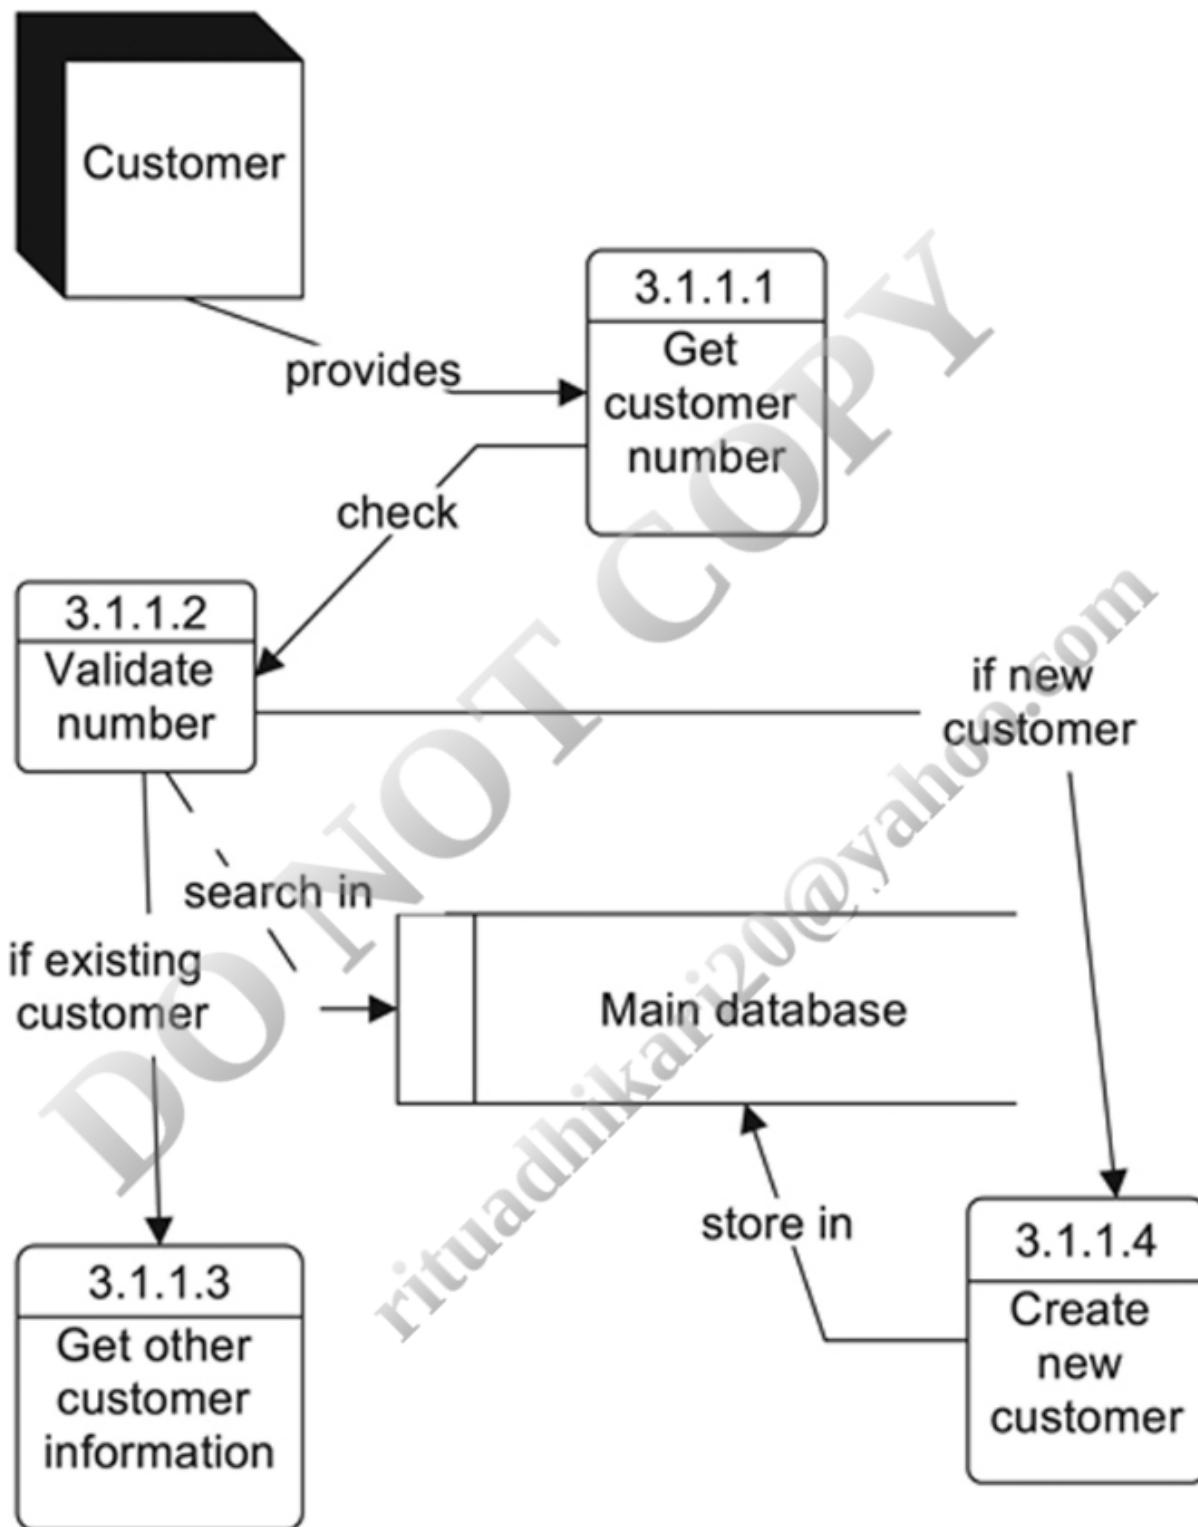


FIGURE 4.20 An explosion of the “Get customer information” process from [Figure 4.19](#).

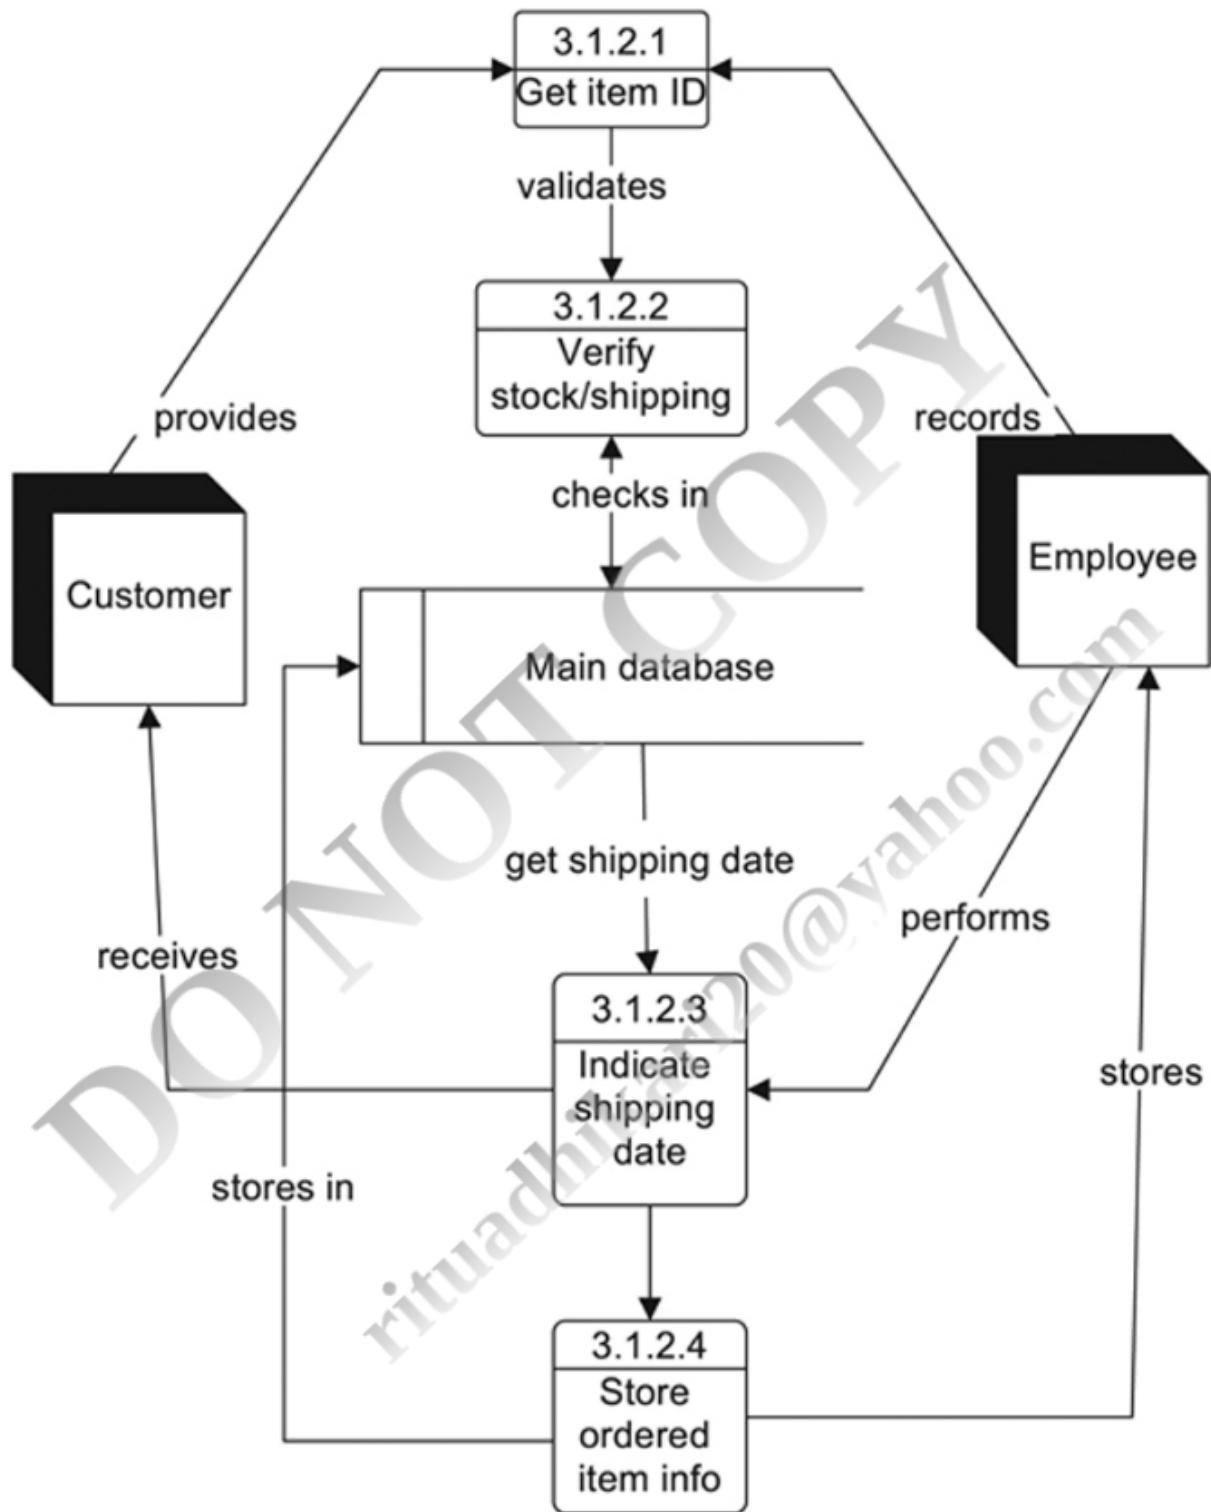


FIGURE 4.21 An explosion of the “Get items ordered” process from [Figure 4.19](#).

Where do the database and the ER diagram fit into all of this? The entire ER diagram is buried inside the “Main database.” In fact, most CASE software

allows you to link your ER diagram to a database's representation on a data flow diagram. Then, you can simply double-click on the database representation to bring the ER diagram into view.

There are a few guidelines you can use to keep data flows and data models separate:

- A data flow shows who uses or handles data. A data model does not.
- A data flow shows how data are gathered (the people or other sources from which they come). A data model does not.
- A data flow shows operations on data (the process through which data are transformed). A data model does not.
- A data model shows how entities are interrelated. A data flow does not.
- A data model shows the attributes that describe data entities. A data flow does not.

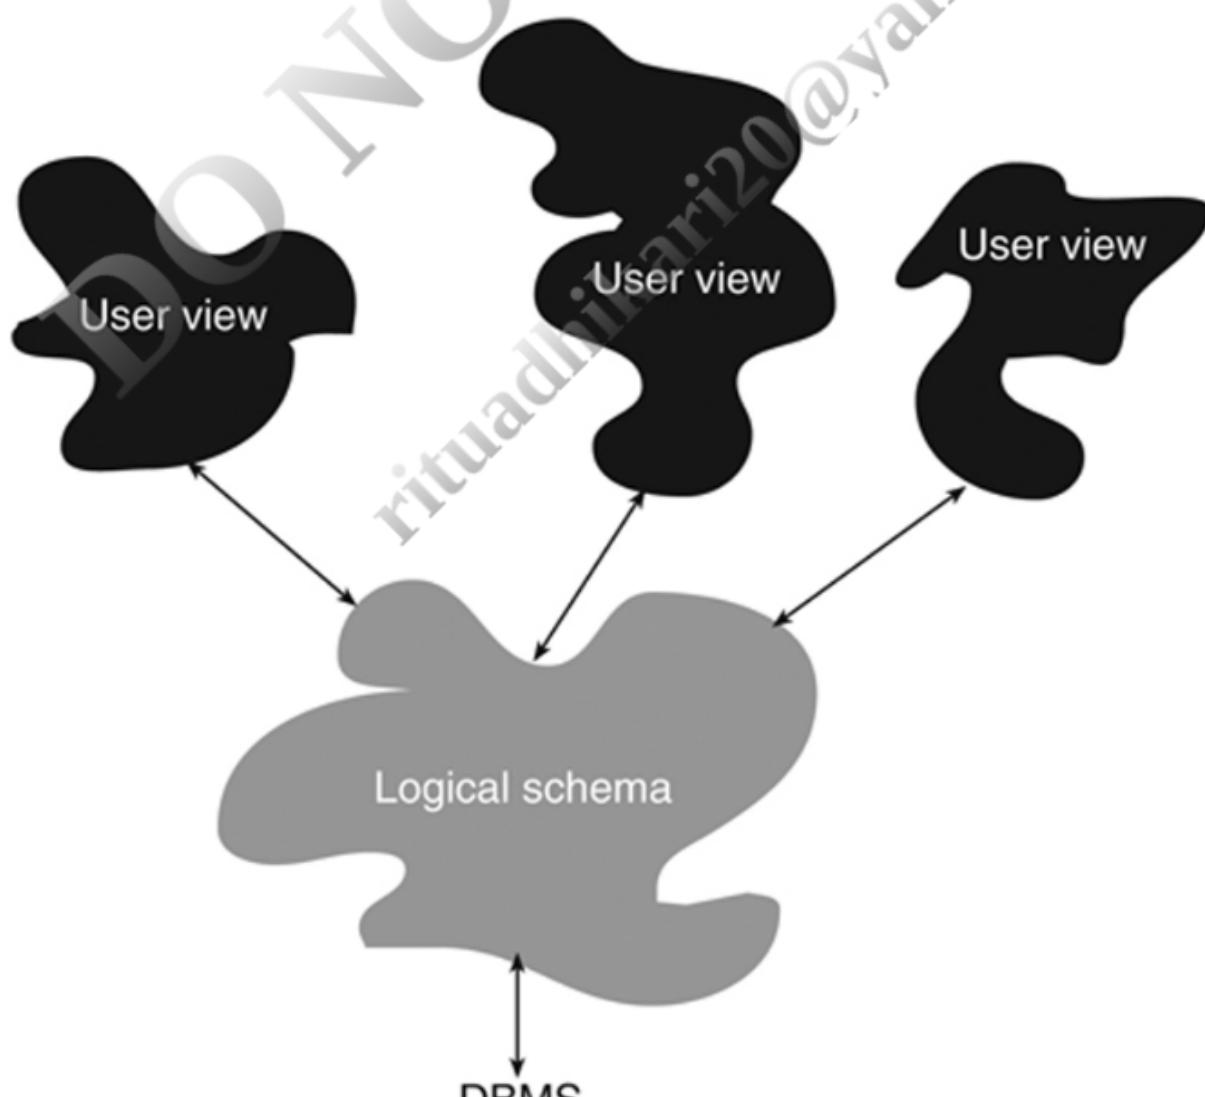
The bottom line is this: A data model contains information about the data being stored in the database (entities, attributes, and entity relationships). If data about an entity are not going to be stored in the database, then that entity should not be part of the data model. For example, although the *Antique Opticals* data flow diagrams show the employee who handles most of the data, no data about employees are going to be stored in the database. Therefore, there is no employee entity in the ER diagram.

Schemas

A completed entity-relationship diagram represents the overall, logical plan of a database. In database terms, it therefore is known as a *schema*. This is the way in which the people responsible for maintaining the database will see the design. However, users (both interactive users and application programs) may work with only a portion of the logical schema. And both the logical schema and the users' views of the data are at the same time distinct from the physical storage.

The underlying physical storage managed by the DBMS is known as the *physical schema*. It is for the most part determined by the DBMS. (Only very large DBMSs give you any control over physical storage.) The beauty of this arrangement is that both database designers and users do not need to be concerned about physical storage, greatly simplifying access to the database, and making it much easier to make changes to both the logical and physical schemas.

Because there are three ways to look at a database, some databases today are said to be based on a *three-schema architecture* (Figure 4.22). Systems programmers, and other people involved with managing physical storage, deal with the physical schema. Most of today's relational DBMSs provide very little control over the file structure used to store database data. However, DBMSs designed to run on mainframes to handle extremely large datasets do allow some tailoring of the layout of internal file storage.



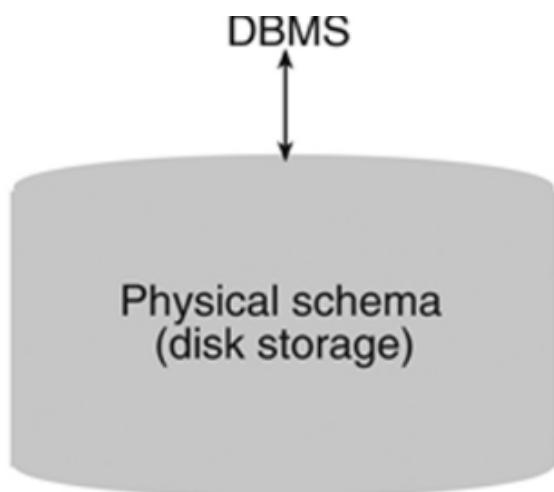


FIGURE 4.22 Three-schema architecture.

Note: DBMSs based on earlier data models were more closely tied to their physical storage than relational DBMSs. Therefore, systems programmers were able to specify physical file structures to a much greater extent. An overview of the older database models can be found in Appendix A.

Data designers, database administrators, and some application programmers are aware of and use the logical schema. End users working interactively, and application programmers who are creating database applications for them, work with user views of the database.

Throughout most of the first portion of this book we will be focusing on the design of the logical schema. You will also learn how to create and use database elements that provide users with limited portions of the database.

For Further Reading

The entity-relationship model was developed by Peter P.S. Chen. If you want to learn more about its early forms and how the model has changed, see the following:

Chen P. The entity-relationship model: toward a unified view of data. *ACM Trans Database Sys.* 1976;1(1):9–36.

Chen, P., 1977. The Entity-Relationship Approach to Logical Database Design. QED Information Sciences, Data Base Monograph Series, No. 6.

Chen, P., 1981. Entity-Relationship Approach to Information Modeling. E-R Institute.

The original work that described the Information Engineering approach can be found in the following:

Martin J. *Information Engineering, Book I: Introduction, Book II: Planning and Analysis, Book III: Design and Construction*. New Jersey: Prentice Hall; 1989.

Finkelstein C. *An introduction to Information Engineering*. Australia: Addison-Wesley; 1989.

For information on UML, see:

Chonoles, M.J., Schardt, J.A., 2003. *UML 2 for Dummies*. For Dummies Press.

Fowler M. *UML Distilled: A Brief Guide to the Standard Object Modeling Language*. third ed. Addison-Wesley Professional; 2003.

Pilone D, Pitman N. *UML 2.0 in a Nutshell*. second ed. O'Reilly Media, Inc; 2005.

For more in-depth coverage of ER modeling, see any of the following:

Baqui S, Earp R. *Database Design Using Entity-Relationship Diagrams*. Auerbach; 2003.

Batini C, Ceri S, Navathe SB. *Conceptual Database Design: An Entity-Relationship Approach*. Addison-Wesley; 1991.

Earp R. *Database Design Using Entity-Relationship Diagrams*. Taylor & Francis; 2007.

Thalheim B. *Entity-Relationship Modeling: Foundations of Database Technology*. Springer; 2000.